

The Case For Language Model Approximated LIKE Predicate

Anonymous Author(s)

ABSTRACT

Modern databases often use the LIKE predicate to search text data. However, when the search condition is interrupted by wildcards, the existing search structure can degrade to a worst-case complexity linear to full table scale, resulting in poor performance. Traditional methods, such as B+-trees, fail to handle wildcards at both ends efficiently. Recent advances in language models offer a promising solution. These models can decode complex LIKE patterns into a small set of candidate values, which are then verified in dataset-size-invariant time via hash table lookups, greatly improving efficiency. However, integrating LLMs into databases faces challenges such as high latency, large storage requirements, and sensitivity to data distribution drifts.

To address these issues, we propose SMILE, a Small Language Model Integrated LIKE Engine that learns column-local character distributions through small but exquisite parameters. Our SMILE acts as a neural translator that converts complex LIKE patterns into their corresponding result sets. Our approach achieves asymptotic complexity improvements while preserving SQL LIKE logic. We conduct comprehensive evaluation across diverse datasets to validate the efficacy of our approach. Our compact SMILE, with a parameter size 5 orders of magnitude smaller than large language models, achieves strong LIKE decoding efficiency and quality. Specifically, SMILE obtains high recall ability while accelerating LIKE by 3 orders of magnitude compared to large language models and sequential scans, 1.8-41.6 times faster than trigram indexes, and 2 orders of magnitude faster than B+-trees. Moreover, our model demonstrates robustness against potential data and query distribution drifts.

CCS CONCEPTS

- Information systems → Query optimization.

KEYWORDS

Query optimization

ACM Reference Format:

Anonymous Author(s). 2026. The Case For Language Model Approximated LIKE Predicate. *Proc. ACM Manag. Data* 4, 1, Article 89 (February 2026), 31 pages. <https://doi.org/10.1145/3786703>

1 INTRODUCTION

Modern database systems heavily rely on the LIKE predicate to efficiently retrieve and compute relevant textual information in relational tuples [14, 63, 75, 80]. This operator is crucial for various applications, including substring matching [66], data cleaning [46, 88, 133], and entity recognition [36].

Please use nonacm option or ACM Engage class to enable CC licenses.
This work is licensed under a Creative Commons Attribution 4.0 International License.
© 2026 Copyright held by the owner/author(s).
ACM 2836-6573/2026/2-ART89
<https://doi.org/10.1145/3786703>



Despite the widespread applications of the LIKE predicate, optimizing it remains a formidable challenge. Traditional B+-tree indexes work well for prefix/suffix patterns [33, 48] but struggle with wildcard-surrounded queries such as "%keyword%" [39, 124]. While trigram indexes support substring searches, they have high storage costs and limited wildcard flexibility [39]. When the queried keyword is interrupted by wildcards, or when the retrieving keyword is a "stop-word" present in all tuples [41], retrieval efficiency drops to $O(N \times \ell)$, where N represents the filtered table size, ℓ is the average string length. Furthermore, trigram indexes are not universally supported across databases [43, 98]. The fundamental limitation of existing approaches lies in their reliance on partial unmasked segments of the LIKE pattern for filtering and search operations, neglecting the holistic context of the query. This constraint significantly undermines their efficacy, particularly when dealing with complex wildcard patterns.

While exact evaluation for arbitrary LIKE patterns inherently requires $O(N \times \ell)$ complexity due to unavoidable full scans for high-cardinality cases (e.g., pure "%"), a critical class of interactive applications requires only a small set of guaranteed matches with sub-second latency. For instance, in data cleaning, in validating Pattern Functional Dependencies (PFDs) [102] such as [Name LIKE "%John%"] → [gender = "M"], within a human-in-the-loop workflow [23, 45, 83]. In this setting, prohibitive scan latencies disrupt the interactive loop, while IR-style approximations are unsuitable as they return logically incorrect matches (e.g., "Joan"). This dilemma motivates an approach rooted in the principles of Approximate Query Processing [24, 100], where trading exhaustive completeness for interactive speed is a well-accepted compromise. We therefore propose a generate-and-validate strategy. If we have a near-oracle predictor to generate a constant-sized candidate set, we can then verify them in $O(\ell)$ time. This eliminates the dependency on dataset size N , delivering the performance expected in interactive systems that present constant-sized result sets [17, 26, 27, 62, 81]. We therefore investigate the feasibility of constructing such an oracle for efficient LIKE query processing.

Fortuitously, Large Language Models (LLMs) [22, 103, 104] approximately satisfy this oracle requirement. Thanks to their pre-trained masked prediction capabilities [37, 90, 115], LLMs can infer missing tokens from incomplete wildcard patterns, making them perfectly suitable for decoding LIKE wildcards into candidates. Figure 1 illustrates the core idea of our approach. For a query such as SELECT * FROM T WHERE Conference LIKE "%IGMO%", we may apply prompt engineering to guide the LLM in understanding that the Conference column contains academic venue names. LLMs can then infer that the pattern "%IGMO%" likely corresponds to the name SIGMOD. We can then verify the presence of SIGMOD via a fast $O(\ell)$ hash lookup. Building on this potential, we proceed to evaluate whether we can significantly reduce the worst-case $O(N \times \ell)$ runtime of such interactive LIKE queries by using language models as efficient autocomplete keyboards—transforming patterns into




Figure 1: Example for language model Approximated LIKE.

a fixed set of candidate strings which are verified in $O(\ell)$ time via hash-based lookups.

However, deploying LLMs to accelerate the LIKE predicate faces a huge obstacle. LLMs are too **large** to fit into conventional DBMS [44, 91]. It therefore brings three barriers: (1) **Latency gap between DB applications and LLMs**: The LLM's response generation incurs second-level delays per query, conflicting with database's demands. For most OLAP scenarios with medium-sized data, a full-table scan (1 second) may be completed before LLMs finish their first sentence (10 seconds). (2) **Storage overheads of LLMs**: The storage space occupied by LLMs is astonishing. For example, DeepSeek-V3, with 671 billion parameters, requires nearly 1TB of storage [120]. This storage consumption can exceed the size of the entire database, significantly reducing the applicability of the method. (3) **Distribution drifts**: In database systems, data evolve continuously, necessitating frequent updates to model parameters [57, 86, 130]. However, fine-tuning LLMs is computationally expensive. To conclude, the "large" nature of existing LLMs presents challenges that hinder the direct deployment of translation LIKE predicates via LLMs. We have to explore whether it is possible to achieve similar LIKE translation functionality with a smaller model.

Fortunately, we observe that the distribution of characters within a single column can be regarded as a unique, special language with limited vocabulary combinations. This observation implies that translating LIKE templates into this column's sub-language can be learned via a small language model with limited parameters, obviating the need for resource-heavy LLMs. This insight motivates us to explore the use of lightweight Small Language Model(SLM) to learn the joint probability distribution of these characters. SLM, with smaller parameter scale, consumes significantly less memory and offers faster inference speed, making it well-suited for latency-sensitive database scenarios. Additionally, the lightweight nature of SLM makes it easier to manage through incremental finetuning in dynamic environments. Therefore, this observation inspires us to investigate whether we can **leverage small-scale language models to learn the language distribution of a single table** while retaining the functionality of LLMs for decoding LIKE predicate.

Based on the above discussions, we propose SMILE (Small Language Model Integrated LIKE Engine), which leverages a lightweight character-level language model to translate LIKE queries into their result sets. This approach tackles the key challenges of LIKE predicate optimization. SMILE learns character distributions within specific column contexts, rather than across broad natural language domains. Coverage in this specific domain needs just a 16MB network, which is 5 orders of magnitude smaller than typical LLMs, yet it still translates LIKE patterns effectively. This small

yet powerful model significantly boosts inference efficiency and reduces maintenance overhead. More importantly, SMILE transforms the inherently difficult LIKE matching problem into a more manageable string candidate verification task. During query optimization, SMILE generates candidate strings that match the LIKE pattern's logical constraints through autoregressive decoding. These candidates are then verified against the database using hash-based lookups. This hybrid approach converts the worst-case $O(N \times \ell)$ full table scans into $O(\ell)$ lookups for valid candidates, achieving substantial complexity improvements while preserving approximate SQL logic. In summary, our work makes the following key contributions:

C1. A Novel Framework for LIKE Predicate Optimization: We introduce a new framework that recasts LIKE predicate optimization as a sequence translation task. This approach employs language models to directly generate candidate tuples from wildcard patterns, fundamentally shifting the execution paradigm from costly linear scans ($O(N \times \ell)$) to an efficient generate-and-validate process.

C2. Lightweight and Scalable SLM Architecture: To overcome the practical limitations of large language models (LLMs), we design and implement SMILE, a lightweight (16MB) and column-specific Small Language Model (SLM). SMILE acts as a specialized neural translator, efficiently learning data distributions to decode LIKE predicates into result sets. We further propose a hybrid strategy that enables our approach to scale to long text columns, which is a critical challenge for existing methods.

C3. Theoretical Framework with Probabilistic Guarantees: We establish a theoretical framework that provides probabilistic approximation guarantees for language-model-based LIKE decoding. By introducing the concept of pattern entropy to quantify query complexity, our framework formally bounds the sampling effort required to retrieve a target number of results with high confidence, establishing a principled link between model capacity, query structure, and performance.

C4. Evaluation Guided by Real-World Workload Analysis: To ensure a realistic evaluation, we first conducted a large-scale survey of 84,047 queries across 30 public databases to characterize real-world LIKE predicate usage. Guided by the identified patterns, we designed a representative benchmark on which SMILE achieves speedups of up to three orders of magnitude over sequential scans, and outperforms specialized trigram and B⁺-tree indexes by 1.8–41.6× and two orders of magnitude, respectively, while using a model five orders of magnitude smaller than mainstream LLMs.

Roadmap: This paper is organized as follows. Section 2 formalizes the SQL LIKE predicate as a sequence-to-sequence task via neural machine translation. Section 3 evaluates state-of-the-art LLMs on LIKE pattern decoding, assessing accuracy, efficiency, and practical viability in database contexts. Section 4 introduces SMILE, a lightweight language model designed for efficient LIKE acceleration. In Section 5, we establish theoretical guarantees and extend the method's applicability. Section 6 presents experimental results comparing SMILE against baselines in accuracy, scalability, and robustness. Section 7 surveys related work, and Section 8 concludes with a summary, limitations, and future directions. Our Appendices provide more details about our collected workloads and experiments [].

233 2 PRELIMINARIES

234 In this section, we first provide an overview of the SQL LIKE predicate
 235 and language models in Section 2.1. Subsequently, we introduce
 236 a novel perspective on the optimization of the LIKE predicate from
 237 the viewpoint of neural language translation in Section 2.2. Section
 238 2.3 explores practical interactive scenarios.
 239

240 2.1 Basic Concepts

241 In this section, we introduce the SQL predicate LIKE and the definition
 242 of language models. Firstly, we delve into the definition of
 243 the LIKE predicate:
 244

245 **The LIKE predicate:** The LIKE predicate facilitates the selection
 246 of tuples within a relational structure where attribute values match
 247 a specified pattern string. Formally, given a relation table T with
 248 attribute A , where the average string length of A is ℓ , and N tuples
 249 in total, the LIKE predicate \mathcal{P} collects all tuples that satisfy \mathcal{P} in re-
 250 sult set $R_1 = \{t \in T \mid t[A] \text{ matches } \mathcal{P}\}$. As standardized in ISO/IEC
 251 9075-2:2023 (Section 8.5, General Rule 3.b.ii) [7], the predicate rec-
 252ognizes two meta characters: (1) The percent symbol (%): Matches
 253 any substring of zero or more characters. (2) The underscore (_):
 254 Matches exactly one arbitrary character.

255 To illustrate the functionality of LIKE predicate, we provide the
 256 following examples. (E1): Pattern LIKE "SIGMOD" yields a condition
 257 where string matches "SIGMOD" string. (E2): Pattern LIKE "SIG%"
 258 matches all strings with the prefix "SIG" through left-anchored pat-
 259 tern matching. (E3): Pattern LIKE "S_GMOD" matches strings length
 260 of six with the first character being "s" and the last being "GMOD".

261 **Problem Definition:** Given a LIKE predicate \mathcal{P} , we propose to
 262 optimize its approximate processing through dual objectives: min-
 263 imizing query latency while maximizing recall of retrieved tuples.
 264 The recall metric is formally defined as: Recall = $\frac{|R_1 \cap R_2|}{|R_1|}$, where R_2
 265 is the retrieved tuples, R_1 is \mathcal{P} 's ground truth result set.

266 Then, we proceed to provide the definition of language models.

267 **Language model:** Given vocabulary $V = \{x_1, x_2, \dots, x_{|V|-2}, \langle SOS \rangle, \langle EOS \rangle\}$ and a sequence of preceding tokens $s_{<t} = (\langle SOS \rangle, x_{\pi 1}, x_{\pi 2}, \dots, x_{\pi t-1})$, the language model \mathcal{M} learns the conditional probability $\Pr(x_{\pi t} \mid x_{\pi 1}, x_{\pi 2}, \dots, x_{\pi t-1})$ and samples the next token $x_{\pi t}$ based on this conditional probability distribution. That is:

$$272 x_{\pi t} \sim \Pr(x_{\pi t} \mid x_{\pi 1}, x_{\pi 2}, \dots, x_{\pi t-1})$$

273 Given a fixed-length string input, the language model performs
 274 autoregressive decoding token-by-token iteratively. This iterative
 275 process begins with the start-of-sequence token $\langle SOS \rangle \in V$. In the
 276 t -th iteration, given the input $s_{<t} = (\langle SOS \rangle, x_{\pi 1}, x_{\pi 2}, \dots, x_{\pi t-1})$,
 277 the language model generates the next token $x_{\pi t}$ based on $s_{<t}$. The
 278 generated token $x_{\pi t}$ is then appended to the historical input to
 279 form $s_{<t+1} = (s_{<t}, x_{\pi t})$. This updated string $s_{<t+1}$ serves as the
 280 input for the next iteration. The iterative process continues until
 281 $x_{\pi t}$ is the end-of-sequence token $\langle EOS \rangle \in V$.

283 2.2 LIKE Predicate Decoding: A 284 Neuro-Linguistic Translation Perspective

285 We reconceptualize LIKE predicate evaluation as a specialized ma-
 286 chine translation task. In this view, a LIKE pattern \mathcal{P} is a source-
 287 language query that must be translated into its logical equivalents
 288 within the database's string domain (the target language). Formally,

let A be the target column domain with N tuples. Conventional
 evaluation computes $\{s \in A \mid s \text{ matches } \mathcal{P}\}$ via deterministic
 approaches and faces a worst-case sequential scan for patterns like
 "%keyword%" in $O(N \times \ell)$ time. In contrast, our approach learns
 a translator $\mathcal{M} : \mathcal{P} \rightarrow \mathcal{A} \subseteq A$, which translates the pattern into
 its results, where:

$$298 \mathcal{A}' = \arg \max_{s \in A} \Pr(s \mid \mathcal{P}) \quad \text{subject to } s \text{ matches } \mathcal{P}$$

300 Here, $\Pr(s \mid \mathcal{P})$ represents the language model \mathcal{M} 's learned dis-
 301 tribution. The verification condition $s \text{ matches } \mathcal{P}$ ensures the cor-
 302 rectness of the result, while the generation process leverages \mathcal{M} 's
 303 logical understanding to bypass exhaustive comparisons. This dual-
 304 phase approach achieves sublinear complexity by restricting pattern
 305 matching to a probabilistically generated candidate set $|\mathcal{A}'| \ll |A|$.

306 This paradigm shift overcomes the fundamental limitations of
 307 traditional indexes through logic-aware holistic decoding. While
 308 trigram indexes fail to fetch wildcard-broken or "stop-word" pat-
 309 terns, and B+-trees stumble on non-prefix alignment, the transla-
 310 tion language model exploits distributional semantics to directly
 311 infer lexically valid matches. Our experimental validation in Sec-
 312 tion 6 demonstrates that this approach maintains SQL-standard
 313 logic while reducing asymptotic complexity from $O(N \times \ell)$ to
 314 $O(\ell) + \mathcal{M}_{decode}$, where \mathcal{M}_{decode} represents the decoding overhead
 315 of modern Transformer architectures.

317 2.3 The Case for Instantaneous, Approximate 318 LIKE Applications

319 Instantaneous feedback on LIKE predicates presents a dilemma
 320 in modern data systems. Traditional query processing guarantees
 321 completeness via exhaustive scans but suffers from high latency.
 322 Conversely, information retrieval techniques—such as full-text in-
 323 dexing [85] or embedding based reranking [58, 74] do not preserve
 324 the logical correctness of SQL LIKE patterns. We argue for a third
 325 paradigm for interactive applications: one that trades completeness
 326 for speed while strictly preserving logical correctness. These ap-
 327 plications need a small, guaranteed-correct subset in interactive
 328 time—a capability current systems lack.

329 **Rule Validation in Interactive Data Cleaning:** Interactive data
 330 cleaning [32, 89, 105] relies on rapid, human-in-the-loop valida-
 331 tion [23, 45, 83] of rules like Pattern Functional Dependencies
 332 (PFDs) [102]. For instance, when validating a rule such as [name
 333 LIKE "%John%"] \rightarrow [gender="M"], an analyst needs immediate
 334 examples to check the logic. A slow full scan disrupts this cognitive
 335 flow, while IR-style approximate matching returns logically irre-
 336 relevant hits (e.g., Joan" or Johnny Appleseed") that obscure logical
 337 correctness. Here, the trade-off is clear: sacrificing completeness
 338 for the interactive speed needed to validate rules is a small price to
 339 pay. Our system delivers a small set of guaranteed-correct matches,
 340 preserving the tight human-in-the-loop process.

341 **Fast Maintenance of Learned Cardinality Estimators:** Learned
 342 cardinality estimators for LIKE predicates [15, 76, 114] must be
 343 continuously retrained on evolving data to mitigate "catastrophic
 344 forgetting" [49, 68]. A cornerstone of this adaptation is experience
 345 replay [107, 145], which necessitates frequent acquisition of ground-
 346 truth labels—i.e., actual query results—for diverse LIKE patterns

from the current dataset. However, full-table scans to obtain these labels incur prohibitive latency, rendering online learning impractical. Our approach provides a remedy by generating a training-batch-sized, logically correct sample in sublinear time. For this task, completeness is unnecessary, but correctness is non-negotiable to avoid label noise, making the trade-off for speed ideal.

Pattern Induction in Information Extraction: Information Extraction (IE) relies on rapid "hypothesize-and-verify" cycles to refine patterns [21, 29, 140]. An engineer refining a pattern LIKE "%John@%.edu%" needs immediate feedback on its matches, but conventional query latency disrupts this iterative process. Our method provides millisecond-level, logically correct responses. Here, the goal is not completeness but to instantly receive a few logical witnesses to the pattern's behavior, allowing developers to validate and refine their logic without delay.

3 LARGE LANGUAGE MODELS: LESSONS LEARNED

In this section, we systematically evaluate the following hypothesis: *Since modern large language models(LLMs) have demonstrated remarkable zero-shot capabilities [40, 95, 131], could they serve as an "off-the-shelf" solution for decoding complex LIKE predicates, bypassing the need for specialized indexes or models?* In this section, we systematically investigate this possibility. We evaluate mainstream LLMs on both decoding quality and efficiency, treating them as a potential baseline. However, our findings reveal fundamental limitations that preclude their direct use: (1) substantial storage requirements, (2) prohibitive inference latency, and (3) a distributional mismatch leading to incomplete answers, which are unacceptable in a database context. These results thus serve as a crucial motivation, demonstrating that a specialized approach is not just an alternative, but a necessity.

Experimental Setup: We evaluate eight state-of-the-art open-source LLMs (ranging from 1.5B to 671B parameters) on LIKE predicate decoding. Our evaluation employs two real-world datasets: *TPC-H (lineitem)* [125], comprising 24M natural-language-like strings, and a recent snapshot of the *IMDB (name)* dataset [61], containing 15M real-world strings from the primaryName column.

To ensure practical relevance, our workload generation methodology is grounded in an analysis of 84,047 real-world LIKE patterns (see Appendix A). We focus on scenarios where conventional prefix/suffix indexes fail—specifically, patterns such as LIKE "%Keyword%" and LIKE "%Key%Key%". We define two primary workloads: **W1 (Syllable-based):** Wildcards mask coherent phonetic syllables, simulating partial recall of word spellings where users reconstruct terms via pronunciation. **W2 (Morpheme-based):** Wildcards obscure morphemes—the smallest meaningful linguistic units—mimicking searches driven by semantic components. To enhance realism and complexity, we randomly replace 0 to 5 characters in the non-wildcard segments of each pattern with underscore ("_") wildcards. This process yields 2,000 test queries. Additional details are provided in Section 6.1.

Decoding Quality Analysis: We measure recall—the fraction of relevant strings correctly generated by each LLM for a given LIKE predicate. As illustrated in Figure 2, decoding accuracy correlates positively with model scale. However, a significant performance




Figure 2: LLM recall performance across datasets and workloads. Left: W1. Right: W2.




Figure 3: Computational cost of LLMs. Left: Inference latency. Right: Model size.

ceiling is evident, as no model approaches perfect recall. Even on the TPC-H dataset, where strings follow easier patterns, the state-of-the-art 671B-parameter model achieves a recall of only approximately 0.8. This persistent gap stems from a fundamental distributional mismatch, or more precisely, a bias mismatch. The open-world knowledge in an LLM's pre-training corpus instills an uncontrolled and undesirable bias towards general natural language. This generalist bias causes the LLM to generate numerous logically plausible strings that are valid in general language but do not exist within the target table's closed-world domain, thereby depressing recall. This issue is exacerbated on the IMDB dataset, whose structured strings deviate more sharply from natural language priors, leading to a more pronounced performance decline. Nevertheless, the core limitation persists even on linguistically coherent data, highlighting a systemic challenge in applying generalist models to specialized data retrieval tasks.

Computational Cost Analysis: Figure 3 quantifies the operational overhead of LLM-based decoding. High-performing models incur substantial latency; notably, DeepSeek-V3—the top performer in recall—exhibits inference times approaching those of a full sequential scan, rendering it unsuitable for latency-sensitive query processing. More critically, storage requirements pose an insurmountable barrier: state-of-the-art LLMs demand gigabytes to terabytes of parameter storage—often exceeding the size of the database itself. This inefficiency arises from the *generalist's burden*: LLMs are trained on diverse tasks (e.g., code generation [47], mathematical reasoning [144], dialogue [28]), resulting in a model whose parameters are overwhelmingly biased towards knowledge domains irrelevant to LIKE decoding. Although LLMs demonstrate zero-shot capability, their immense computational and storage demands are fundamentally incompatible with the efficiency constraints of database systems, where every CPU cycle and byte is carefully optimized.

Key Insights: 1. *Foundational Decoding Capability:* LLMs can map wildcard patterns to logically plausible completions,

achieving high recall (up to 80%) on linguistically coherent data by leveraging learned language distributions.

2. The Prohibitive Cost of Scale: High accuracy requires massive models, incurring severe latency and storage overheads—a “success tax” that renders direct deployment economically and operationally infeasible.

3. The Generalist’s Bias: General-purpose knowledge introduces an uncontrolled and inefficient bias; most model capacity is consequently irrelevant to LIKE decoding, violating core database design principles of specialization and resource frugality.

These results offer a crucial insight. The initial enthusiasm for general-purpose LLMs is met by the harsh reality of the “generalist’s burden”—their vast knowledge introduces uncontrolled bias and prohibitive overhead, making them impractical for production database use. This realization motivates a deliberate shift in design philosophy: we argue that a workload-aware, specialized solution is a more principled and efficient engineering choice. Consequently, we propose a lightweight small language model designed to overcome the exact limitations identified in this section.

4 SMALL LANGUAGE MODELS: AN ECONOMICALLY VIABLE ALTERNATIVE

To overcome the limitations imposed by the LLMs’ generalist bias, we propose SMILE (Small language Model Integrated LIKE Engine), a cost-effective solution that leverages SLMs to enhance the efficiency of LIKE predicate processing. SMILE is designed around a core principle: the intentional learning of a potent *inductive bias*. Instead of relying on open-world knowledge, SMILE learns the specific character-level distribution of a target column. This learned domain prior is the key to its efficiency and high recall, enabling it to act as a specialized neural translator that operates within the closed world of the database column.

4.1 SMILE Overview

As depicted in Figure 4, SMILE operates through three core stages: offline preparation of model and training workloads, online small language model driven translation of LIKE predicates, and hash-based validation of candidate strings. SMILE aims to optimize low-cardinality LIKE query patterns where dataset-size-independent decoding acceleration via SLM becomes feasible. This section provides a concise workflow overview. We formally introduce SMILE’s SLM architecture and decoding strategies in Sections 4.2 and 4.3, respectively. Section 4.4 discusses decoding optimization techniques, while Section 4.5 details our query routing mechanism for identifying targetable low-cardinality queries.

1. Offline preparation: We initially prepare the training workloads offline. Subsequently, we train a classifier on the target column A and a character-level small language model using the prepared workloads. Inspired by prior works on learned database [111, 132], we first cluster historical queries into high-cardinality and low-cardinality types. We then train a binary classifier that classifies query patterns into these two categories. For low-cardinality patterns, we statistically model their wildcard distributions (“%” and “_” placements) to generate synthetic training data. We firstly sample N_s strings from A and obtain training sampling set $\mathcal{S} = \{s \mid s \in A\}, |\mathcal{S}| = N_s$. Then, we inject wildcards into the sampling set \mathcal{S} to

form supervised training information $\langle \mathcal{P}_i, s_i \rangle$ pairs where wildcard pattern \mathcal{P}_i constitutes a source language and original tuples s_i form the target language. We employ teacher forcing with cross-entropy loss to train SMILE’s SLM:

$$\mathcal{L} = - \sum_{\langle \mathcal{P}_i, s_i \rangle} \sum_{x_{ij} \in s_i} x_{ij} \log \mathcal{M}(\mathcal{P}_i)_j$$

Here, x_{ij} represents the true distribution of target string s_i ’s j -th token, and $\mathcal{M}(\mathcal{P}_i)_j$ represents the language model’s output probability distribution for the j -th output token given wildcard input \mathcal{P}_i . The training approach mentioned above can be adapted for data updates. In scenarios where data undergoes periodic changes, we can similarly generate training samples and wildcard patterns from the new data distribution to fine-tune the model parameters. This training process makes SMILE acquires its inductive bias. By learning exclusively from the target column’s data, the model internalizes a domain-specific prior that constrains its predictive distribution to the statistical patterns of this closed world, effectively eliminating the out-of-domain generation that plagues generalist models.

2. Online SLM translation: For unknown queries, we first utilize a character-level language classifier to route testing query workloads. Subsequently, based on the classification outcomes, we meticulously select wildcard patterns with low-cardinality from the testing workloads. These low-cardinality wildcard patterns can be efficiently translated into candidate sets in $O(\ell)$ time. For these low-cardinality queries, we employ a neural character Transformer to translate the source LIKE patterns into target strings in the database. However, if the query has high-cardinality (e.g., pure “LIKE %”), even if language model can successfully decode such queries, the number of samples generated by the language model is proportional to the scale of the entire database, which restricts the space for the language-model-based optimization. Therefore, we directly deliver these query patterns for the database to execute.

3. Hash verification: After we obtained the constant scale candidate set via the generation of our SLM, we verified our candidate set using hash index, allowing results to be checked in $O(\ell)$ time.

4.2 Transformer: The Efficient LIKE Translator

We find the Transformer architecture [16, 64, 127], originally developed for machine translation, is exceptionally well-suited for our LIKE predicate decoding task. This suitability stems from two of its inherent properties:

(A1). Intrinsic Seq2Seq Architecture: The “%” wildcard implies that the length of a matching string is independent of the pattern’s length. This input-output length mismatch renders fixed-size architectures, such as standard MLPs, unsuitable for learning this mapping [16, 118]. The Transformer’s encoder-decoder structure naturally addresses this challenge. The encoder maps the variable-length input pattern to a latent representation $\mathbf{H} \in \mathbb{R}^{m \times d}$. The decoder then autoregressively generates the output sequence y_t conditioned on \mathbf{H} , terminating upon generating an $\langle \text{EOS} \rangle$ token. This mechanism effectively decouples the input and output lengths, making it ideal for the variable-length translation required by LIKE predicates.

(A2). Dynamic Attention Mechanism: When decoding a LIKE predicate, different segments of the input pattern have varying importance for generating the output string. The attention mechanism




Figure 4: Overview on SMILE.

enables the model to dynamically focus on the most salient input tokens for each decoding step by adaptively weighting the input sequence. Formally, given an input sequence of d -dimensional embeddings $\mathbf{X} = (\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m)$, attention computes Query, Key, and Value matrices via learnable linear projections: $\mathbf{Q} = \mathbf{X}\mathbf{W}_Q$, $\mathbf{K} = \mathbf{X}\mathbf{W}_K$, $\mathbf{V} = \mathbf{X}\mathbf{W}_V$, with projection matrices $\mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V \in \mathbb{R}^{d \times d}$. The rows of \mathbf{Q} , \mathbf{K} , and \mathbf{V} represent the query, key, and value vectors, respectively. The core of the mechanism lies in computing relevance scores:

$$\alpha_{i,j} = \frac{\exp(\mathbf{q}_i \mathbf{k}_j^\top / \sqrt{d})}{\sum_{k=1}^m \exp(\mathbf{q}_i \mathbf{k}_k^\top / \sqrt{d})}$$

where the scaled dot-product $\mathbf{q}_i \mathbf{k}_j^\top / \sqrt{d}$ measures feature compatibility. These normalized coefficients implement a soft reranking mechanism, where the attention coefficients assess the importance of the current feature contribution to LIKE decoding. Based on the attention weights, the context matrix $\mathbf{Z} \in \mathbb{R}^{m \times d}$ is computed via value aggregation as follows, where \mathbf{z}_i denotes the i -th row of \mathbf{Z} :

$$\mathbf{z}_i = \sum_{j=1}^m \alpha_{i,j} \mathbf{v}_j \quad \text{for } i = 1, \dots, m,$$

This reweighting allows automatic emphasis on logically critical subpatterns (non-wildcard segments) while maintaining long-range dependency awareness - crucial for reconstructing masked segments between wildcards. For example, given a LIKE pattern "SIG_O_" to decode, we need to pay more attention to "SIG" than "O" as the prefix "Special Interest Group"("SIG") gives us more details of the underlying matched result than "O". Therefore, these architectural advantages make Transformers naturally suitable for LIKE predicate decoding: The dynamic length adaptation handles

arbitrary wildcard expansions, while the attention mechanism automatically focuses on discriminative subpatterns to efficiently decode string patterns.

4.3 Capture LIKE pattern in a Seq2Seq manner

SMILE adapts Transformer to perform character-level translation of LIKE predicates. As illustrated in Figure 4, the model processes a LIKE pattern through distinct encoding and decoding stages.

Encoding Stage: The encoding stage transforms an input LIKE pattern \mathcal{P} of length m into a dense vector representation $\mathbf{H} \in \mathbb{R}^{m \times d}$ via three sequential operations: (1) **Character-level Embedding**: Each character in the input pattern is mapped to a d -dimensional vector using a trainable embedding matrix $\mathbf{Emb} \in \mathbb{R}^{|V| \times d}$, where V is a column-specific vocabulary. (2) **Positional Encoding**: To inform the model of the character sequence order, which is otherwise ignored by the self-attention mechanism, we inject positional information. This is achieved by summing the character embeddings with a standard sinusoidal positional encoding [127]:

$$PE_{(pos,2i)} = \sin(pos \times 10000^{-\frac{2i}{d}}), \\ PE_{(pos,2i+1)} = \cos(pos \times 10000^{-\frac{2i}{d}}),$$

where pos is the character's position and i is the dimension index ($i \leq d$). This parameter-free approach facilitates generalization to variable-length patterns, which is crucial for robustly handling wildcards. (3) **Self-attention Encoding**: Finally, a Transformer encoder layer processes the resulting sequence. It employs a multi-head self-attention mechanism, followed by layer normalization and a residual connection, to capture long-range dependencies between characters in the pattern.

Decoding stage: The decoder learns a scoring function that maps the encoded LIKE pattern representation $\mathbf{H} \in \mathbb{R}^{m \times d}$ and the currently generated sequence $s_{<t} = (\langle SOS \rangle, x_{\pi 1}, x_{\pi 2}, \dots, x_{\pi t-1})$ to

Algorithm 1 Temperature annealed decoding for LIKE patterns

Input: Encoded LIKE pattern $\mathbf{H} \in \mathbb{R}^{m \times d}$, temperature T , beam size K , decode length ℓ , Hash Verifier *Hash*

Output: Valid candidates $C_{\text{valid}} \subseteq C$

```

1: procedure DECODELIKE( $\mathbf{H}, T, K, \ell, \text{Hash}$ )
2:    $\mathbf{H}_{\text{tile}} \leftarrow \text{repeat}(\mathbf{H}, K)$             $\triangleright$  Tile pattern representation
3:    $\mathbf{s}_t \leftarrow [\langle \text{SOS} \rangle]^{\times K}; \text{done} \leftarrow \text{False}^K;$        $\triangleright$  Init sequences
4:   for  $t = 1$  to  $\ell$  do
5:      $D_t \leftarrow \text{Decoder}(\mathbf{H}_{\text{tile}}, \mathbf{s}_t)$            $\triangleright$  Decode
6:      $\mathbf{P}_t \leftarrow \frac{\exp(D_t/T)}{\sum_{\epsilon \leq |V|} \exp(D_t[\epsilon]/T)}$      $\triangleright$  Temperature scaling
7:      $y_t \leftarrow \text{sample}(\mathbf{P}_t);$ 
8:      $\mathbf{s}_t \leftarrow \mathbf{s}_t + (y_t \cdot \neg \text{done}); \text{done} \leftarrow \text{done} \vee (y_t == \langle \text{EOS} \rangle)$ 
9:     if all(done) then break
10:     $C_{\text{valid}} \leftarrow \{s_t \in \mathbf{s}_t \mid \text{Hash}[s_t] = 1\}$        $\triangleright$  Hash verification
11:   return  $C_{\text{valid}}$ 

```

the vocabulary score distribution D_t of the t -th token. This function is implemented via a Transformer-based architecture, maintaining dimensionality consistency across layers. The decoder then projects these hidden states to the score vector $D_t \in \mathbb{R}^{|V|}$:

$$D_t = \mathbf{W}_o \mathbf{h}_t^N + \mathbf{b}_o$$

where $\mathbf{h}_t^N \in \mathbb{R}^d$ denotes the final layer's hidden representation at step t , $\mathbf{W}_o \in \mathbb{R}^{|V| \times d}$ is the output weight matrix, and $\mathbf{b}_o \in \mathbb{R}^{|V|}$ is the bias vector with $|V|$ being the vocabulary size. The resulting weight vector from the decoder contains the likelihood scores for all tokens in the vocabulary. The subsequent section discusses an efficient decoding strategy based on these scores.

4.4 Translating LIKE Via Temperature Annealed Importance Sampling

Translating a LIKE predicate that matches multiple strings requires generating a diverse set of candidates. Standard greedy decoding, however, often yields repetitive outputs, as it consistently favors high-probability tokens and fails to explore the full solution space [146].

To address this, we employ a temperature-annealed sampling strategy, a technique adapted from large language model decoding [20, 22, 37, 146] to balance exploration and exploitation. Given a LIKE pattern \mathcal{P} over column A with vocabulary V , SMILE's decoder computes a score distribution $D_t \in \mathbb{R}^{|V|}$ at each step t , conditioned on \mathcal{P} and the prefix $s_{<t}$. We then apply temperature scaling to this distribution to promote diversity:

$$\Pr(x_{\pi t} = x_k) = \frac{\exp(D_t[k]/T)}{\sum_{\epsilon \leq |V|} \exp(D_t[\epsilon]/T)}$$

Algorithm 1 outlines this process. To generate K candidates concurrently, the encoded LIKE pattern \mathbf{H} is tiled, and each sequence is initialized with a $\langle \text{SOS} \rangle$ token (Lines 2-3). The algorithm then iteratively samples tokens. In each step, the decoder generates logits D_t , which are scaled by temperature T to form a probability distribution \mathbf{P}_t (Line 5). A new token y_t is sampled from \mathbf{P}_t and appended to each active sequence (Lines 6-7). The process terminates once all sequences have generated an $\langle \text{EOS} \rangle$ token. Finally, a hash-based verifier validates each candidate string in $O(\ell)$ time (Line 10).

The overall time complexity is $O(D \cdot \ell)$, where D is the decoder's computational cost and ℓ is the decode length. This complexity is independent of the database size N , ensuring scalability.

4.5 Cardinality-Aware Classification

While SMILE achieves asymptotic acceleration for low-cardinality LIKE queries, high-cardinality patterns (e.g., pure LIKE "%") pose challenges. To decode type LIKE "%" predicate would require the sampling budget being linear to the table scale. Therefore, such exhaustive generation for non-selective patterns approaches $O(N \times \ell)$ complexity, negating our complexity improvements.

To resolve this tension, we introduce a cardinality-aware classifier that dynamically routes queries between the SLM translator and traditional execution methods. The classifier predicts whether a given LIKE pattern will exceed predefined cardinality thresholds proportional to the sampling budget, and we reuse the neural translator's pre-trained encoder (described in §4.3) to obtain $m \times d$ dimensional input patterns' hidden representation \mathbf{H} . We then use average pooling to map \mathbf{H} into d dimensional summarization vectors. This summarization captures both syntactic features (wildcard positions, pattern length) and distribution characteristics (expected character distributions). A two-layer MLP with ReLU activation then maps these representations to cardinality class probabilities:

$$\Pr(\text{high-card}) = \sigma(\mathbf{W}_2 \cdot \text{ReLU}(\mathbf{W}_1 \cdot \text{Avg}(\mathbf{H}) + \mathbf{b}_1) + \mathbf{b}_2)$$

where $\text{Avg}(\mathbf{H})$ denotes average pooling of the input LIKE pattern's encoded representation \mathbf{H} , and σ represents the sigmoid function. The model trains on query logs annotated with true cardinality values, optimizing binary cross-entropy loss. In practice, this lightweight classifier (6MB parameters) introduces negligible overhead (0.1ms/query) while achieving accurate query routing in our experiments, enabling reliable fallback to full scans for low-selectivity patterns without compromising low-latency acceleration for targeted queries.

5 GUARANTEES AND EXTENSIONS

Having established the architecture of SMILE, we now turn to its theoretical underpinnings and practical extensibility. This section first develops a framework for probabilistic approximation guarantees, formally quantifying the relationship between the sampling effort of our SLM-based translator and the completeness of the query result. Building upon this theoretical foundation, we then explore practical extensions of our methodology.

5.1 Probabilistic Guarantees for Approximation

To develop a principled understanding of our language model based approach for LIKE predicate approximation, we construct a theoretical model. The goal of this model is to quantify the relationship between a query's complexity, the model's capacity, and the number of samples, n , required to retrieve at least K matching strings with confidence $1 - \delta$.

Our framework's central premise is to contrast the language model's objective with the conventional LIKE execution, which is a *discriminative* process: given a fixed, finite set of stored strings, the system identifies which ones satisfy the LIKE pattern. In contrast, our language model tackles this a *generative* task: it must infer the syntactic constraints implicit in the pattern and synthesize valid

strings *de novo*. Hence, the difficulty of this task depends not merely on the empirical data distribution of any specific database instance, but more on the intrinsic structural complexity of the pattern itself. **Assumption 1 (Task Complexity as Generative Entropy).** We model the complexity of a LIKE pattern \mathcal{P} by the difficulty of the generative task it represents. This difficulty is proportional to the size of its corresponding language, $L(\mathcal{P})$ —the space of all possible valid strings. To formally quantify this *intrinsic generative complexity* while abstracting away from the contingent data distribution of any particular database, we employ the principle of maximum entropy. This leads us to assume a uniform distribution over all syntactically valid strings in $L(\mathcal{P})$, thereby isolating the complexity inherent in the pattern’s structure.

Definition 1 (Pattern Entropy). Under this generative framework, for a LIKE pattern \mathcal{P} over an alphabet Σ , we define the **pattern entropy**, $H(\mathcal{P})$, as the Shannon entropy of the uniform distribution over its language $L(\mathcal{P})$, measured in nats:

$$H(\mathcal{P}) = \ln |L(\mathcal{P})| \quad (1)$$

This value serves as a proxy for the syntactic learning difficulty posed to the generative model.

THEOREM 5.1 (ENTROPY OF A LIKE PATTERN). Given a LIKE pattern of length x with p underscore ("_) and q percent ("%") wildcards over alphabet Σ , no two percent signs adjacent, and a maximum generated-string length of m , the pattern entropy $H(\mathcal{P})$ is:

$$H(\mathcal{P}) = p \ln |\Sigma| + \ln \left(\sum_{s=0}^{m-x+q} \binom{s+q-1}{q-1} |\Sigma|^s \right) \quad (2)$$

The proof is detailed in Appendix I.1. Theorem 5.1 provides a closed-form method to compute this complexity measure. The total entropy of a workload \mathbb{W} is the sum of the entropies of its constituent patterns: $H(\mathbb{W}) = \sum_{\mathcal{P} \in \mathbb{W}} H(\mathcal{P})$. Next, we model the language model’s capabilities in a simplified manner.

Assumption 2 (Model Information Capacity). We followed Meta AI’s recent analyses [97], postulate that a model with $|\theta|$ parameters has an effective information capacity, M , that is proportional to its size, in a linear relationship:

$$M = \alpha |\theta| \quad (3)$$

where α is an architecture-dependent constant. Within our framework, M represents an abstracted upper bound on the total task complexity $H(\mathbb{W})$ that the model can reliably internalize.

Assumption 3 (Capacity-Limited Performance). We postulate a relationship between the model’s capacity and its generative success probability, p . We conceptualize this as a resource allocation problem. If the model’s capacity M meets or exceeds the task’s complexity $H(\mathbb{W})$, its performance is primarily limited by other factors (e.g., training quality, decoding stochasticity), which we bound with a universal constant $c \in (0, 1]$. If the model is under-capacitated ($M < H(\mathbb{W})$), we assume its performance degrades proportionally to the ratio of available capacity to task demand. This leads to the following lower bound on the success probability:

$$p \geq c \cdot \min \left(1, \frac{M}{H(\mathbb{W})} \right) = c \cdot \min \left(1, \frac{\alpha |\theta|}{H(\mathbb{W})} \right) \quad (4)$$

This assumption provides a tractable, albeit simplified, link between model scale and task complexity, enabling the derivation of sampling bounds.

THEOREM 5.2 (SAMPLING COMPLEXITY FOR LIKE APPROXIMATION). Under the assumptions of our theoretical model, for a query workload with entropy $H(\mathbb{W})$, to retrieve at least K correct results with confidence $1 - \delta$, the required number of samples n from a model with $|\theta|$ parameters must satisfy:

$$n \geq \begin{cases} \frac{1}{c} \left(K + \ln(1/\delta) + \sqrt{(2K + \ln(1/\delta)) \ln(1/\delta)} \right) & \text{if } H(\mathbb{W}) \leq \alpha |\theta| \\ \frac{H(\mathbb{W})}{c \cdot \alpha |\theta|} \left(K + \ln(1/\delta) + \sqrt{(2K + \ln(1/\delta)) \ln(1/\delta)} \right) & \text{otherwise} \end{cases} \quad (5)$$

The proof is detailed in Appendix I.2. Theorem 5.2 is the central result of our model, establishing a quantitative link between a workload’s intrinsic complexity ($H(\mathbb{W})$), the model’s abstracted capacity ($|\theta|, \alpha$), and the sampling budget (n) required for a given performance target (K, δ). Crucially, the generative viewpoint of our model implies that for a consistent query distribution, the syntactic rules learned from one data sample can generalize to entirely disjoint data partitions. This property is instrumental for handling evolving databases and enables rapid model deployment, a claim we validate empirically in Section 6.3 and 6.4.

5.2 Discussions

We discuss SMILE’s extensibility to long text columns and its broader applications.

Scalability to Long Text Columns. Deep sequence models for LIKE optimization struggle with long text columns [15, 76, 114]¹, a limitation also present in our baseline full-string generation approach. The core challenges are twofold: (1) long text implies longer m in Theorem 5.1, thus greater entropy, impedes learning the data distribution for exact matches, and (2) the linear cost of autoregressive inference becomes prohibitive. To address this, we propose a hybrid strategy that reframes SMILE as a prefix-guided query rewriter. Instead of generating a full string, the SLM produces a fixed-length, high-selectivity prefix. This prefix is then used to probe a standard index (e.g., a B⁺-tree), converting a LIKE "%keyword%" query into an efficient LIKE "prefix%" lookup. This design solves both issues: it lowers the target entropy ($H(x_{\text{prefix}}) < H(x_{\text{full}})$), improving sampling success per Theorem 5.2, and decouples inference cost from string length to ensure scalability. We believe this is the first practical deep learning approach for LIKE acceleration that scales to the maximum wildcard lengths (e.g., 8KB) supported by modern database systems.²

Generalization to Other Tasks. SMILE’s core methodology is broadly applicable to other tasks. For instance, SMILE can accelerate approximate regular expression (regex) matching. By modeling regex as a sequence-to-sequence task, it can significantly improve performance in interactive data exploration and cleaning workflows. The approach also extends to optimizing fuzzy string joins (e.g., "Orders.CustomerName" with "Customers.Name"). A trained SLM

¹Recent deep LIKE CardEst works use strings of 50–100 length in average, none exceeding 500. [14, 75, 76]

²For example, 255 bytes in SAP [110], 4KB in DB2 [60], and 8KB in SQL Server [93].

can act as a candidate generator, mapping an input like "J. Smith" to high-probability matches such as "John Smith". This transforms an expensive fuzzy join into a series of efficient, model-guided lookups. While primarily designed for single-column approximate queries, SMILE can be readily extended to optimize complex analytical queries, serving as a specialized form of Approximate Query Processing [24, 100]. We employ SMILE as a query rewriter for queries with non-sargable "LIKE" predicates: it generates a candidate set of matching strings and transforms the predicate into an equivalent "WHERE column IN (...) clause. This simple yet powerful rewrite makes the predicate index-friendly, enabling the database optimizer to leverage efficient hash-based lookups instead of costly scans. If the cardinality estimator in the database predicts high selectivity for this predicate, SMILE can be enabled to perform the above rewriting. We validated this strategy end-to-end on relevant queries from TPC-H and IMDB-JOB benchmarks in Sec 6.2 and Appendix H.

6 EXPERIMENTS


In this section, we provide an experimental analysis of SMILE in accelerating approximated LIKE pattern. We use our experiment results to answer the following questions:

(Q1) Effectiveness: Can the proposed language-model-based LIKE predicate decoding strategy significantly reduce the execution time and cost of LIKE predicates compared to traditional methods and state-of-the-art large language models across multiple datasets? Can it generate sufficiently high-quality candidate sets to recall potential query results? Can it play a useful optimization role in practical complex workloads? (Section 6.2)

(Q2) Scalability: Is the method's scalability in training and inference time sufficiently effective? Is the method's space occupancy sufficiently low? How does SMILE's performance and accuracy scale with increasing workload entropy and string length, particularly on billion-record and long-text datasets? (Section 6.3)

(Q3) Robustness: If data distribution drifts occur, can the method be updated in time to adapt to data changes? If the query template distribution changes, can the method still effectively support unknown query template estimation without retraining? (Section 6.4)

6.1 Experimental Setup



(a) Baseline Configuration Tuning(Partition=32).

Hardware Configuration		Hourly Rental Cost (\$)
1 × AMD EPYC 9654 CPU (96 Cores)	[30]	\$1.50
1 × AMD EPYC 9654 CPU (16 Cores)	[30]	\$0.25
1 × NVIDIA H20 GPU + 1 × AMD EPYC 9654 CPU (16 Cores)	[13]	\$0.98

Figure 5: Configuration analysis and hardware cost.

Datasets: We evaluate LIKE predicate performance on six diverse datasets, with detailed statistics summarized in Table 1. Four are standard benchmarks in string query research [14, 75, 137]: (1) **TPC-H** [125], comprising 24 million rows from the comment column of

the lineitem table (SF=4); (2) **IMDB** [61], with 15 million entries from the primaryName column of the Name table of the **real-world** IMDB dataset; (3) **Reddit** [96], containing 2 million **real-world** reddit usernames; and (4) **WIKI** [129], consisting of 4 million **real-world** Wikipedia article titles. To further assess scalability and robustness, we use two large-scale, **real-world** datasets: (5) **Red-Pajama** [123], a set of 1 billion unique string windows (max length 40) extracted from a 1TB corpus; and (6) **News Room** [56], used for long-text evaluation with 1.2 million articles averaging 3962 characters in length.

Table 1: Statistics for the datasets used in our evaluation.

	Reddit	WIKI	IMDB	TPC-H	Redpajama	NewsRoom
Total	2,000,000	4,000,000	14,757,614	23,996,604	1,000,000,000	1,212,741
Unique	1,999,847	4,000,000	11,278,588	15,189,929	1,000,000,000	1,200,762
Min Length	1	1	1	8	1	27
Max Length	20	257	105	43	40	18386
Avg Length	11.00	19.91	13.51	26.23	36.61	3962.47

Workload Design. To evaluate the robustness of our method, we designed four LIKE workloads (W1-W4) that span from realistic to challenging scenarios. **Realism-Driven Workloads (W1 & W2).** We derived W1 and W2 from an analysis in Appendix A based on our collected 84,047 production LIKE queries(available at [8]). The analysis revealed that challenging infix searches (e.g., %Keyword%, %Key1%Key2%) are prevalent. We found that in such queries, the substrings masked by wildcards often correspond to linguistic units like morphemes or syllables, rather than arbitrary characters. Accordingly, we generate realistic patterns by splitting string into list of linguistic units using [9, 10] and replacing 0-2 consecutive linguistic units with a single "%" wildcard. This method covers over 50% of the production cases (Appendix A.3). We therefore define two specific workloads: **W1 (Syllable-based):** Simulates phonetic recall by masking syllables, the basic units of pronunciation [108]. For example, "SIGMOD CONFERENCE" can yield the pattern %MOD CON%EN%. **W2 (Morpheme-based):** Models semantic-driven searches by masking morphemes, the smallest meaningful units in a language [42, 50]. This simulates users searching for concepts embedded within larger words. The detailed construction algorithm of these workload is described in Appendix E. **Benchmark Workloads (W3 & W4).** To benchmark against established standards and test performance boundaries, we include two additional workloads: **W3 (CLIQUE Workload):** Adopts the standard suite of LIKE patterns from CLIQUE [75] (e.g., %Key, Key%, %Key%Key%) to evaluate performance on common predicate forms. **W4 (LPLM Workload):** Uses the generator from LPLM [14] to create challenging patterns with high wildcard density and low semantic content (e.g., \$0%\$1%\$2%...\$n%).

Finally, our analysis in Appendix A.4 and A.6 shows that single-character "_" wildcards are statistically independent of "%" wildcards, and usually appear 0-5 times per query, with occurrences of 6-7 times being less than 0.1%. Based on this finding, we augment all generated patterns across W1-W4 by randomly replacing 0-5 characters in each key with a "_", enhancing the realism of the final query predicates.

Baseline competitors: We select the following baseline competitors as they have proven their effectiveness in realistic database LIKE pattern matching deployments of PostgreSQL, or achieved

state-of-the-art open-sourced text pattern understanding performance in relevant NLP tasks [34, 120, 121]. The baselines include:

1. *Prefix B+-tree*: Accelerates prefix LIKE patterns through prefix-optimized B+-tree structure.
2. *Suffix B+-tree*: Accelerates suffix LIKE patterns using reverse-ordered B+-tree structure.
3. *Prefix+Suffix B+-tree*: Accelerates suffix LIKE patterns using combined prefix and suffix B+-trees.
4. *Sequential scan*: Baseline full-table scanning approach.
5. *GIST index* [39]: Accelerates LIKE via Generalized Search Tree.
6. *GIN index* [101]: Accelerates LIKE through Generalized Inverted Index.
7. *GIN index (Approximate)* [39]: Approximate LIKE matching using inverted index, following a similarity filtering technique in [25, 55].
8. *Qwen2.5-7B* [121]: 7B-parameter LLM developed by Qwen-Team.
9. *Qwen2.5-32B* [121]: 32B-parameter LLM developed by Qwen-Team.
10. *DeepSeek-V2.5* [34]: 236B-parameter LLM developed by Deepseek.
11. *DeepSeek-V3* [120]: State-of-the-art open-sourced 671B LLM.

Experimental settings: The language model of SMILE is configured with a hidden layer size of 512 and a single attention layer. During training, we set the learning rate to 3×10^{-4} and the batch size to 16384. For the SMILE's classifier, we maintain the same hidden layer size of 512 but adjust the batch size to 2048 and the learning rate to 1×10^{-4} . We establish a classification threshold of 16 for the query workload cardinality. When training the small language model, we utilize 200 sampled batches to ensure effective learning and model convergence. During inference, we set the sampling number being 64. Our SMILE is trained and evaluated on a server with a Nvidia H20 GPU and a 96-core AMD EPYC 9654 CPU. The PostgreSQL version is 16.3. To ensure a fair comparison, we first tune the optimal configuration for baselines via W3 of the TPC-H dataset on the 96-core CPU. As shown in Figure 5(a), their performance saturates at 16 parallel workers due to Amdahl's Law [12]. Beyond this point, more cores offer no advantage. Consequently, we configured all CPU baselines with 16 workers and 32 partitions for the main experiments. Our 96-core server's resources can be partitioned, simulating a multi-tenant environment. Therefore, to provide a practical and equitable comparison of economic efficiency, our cost analysis for all methods is normalized to the price of a 16-core CPU setting, which represents the most cost-effective hardware configuration. This methodology ensures a fair comparison of economic efficiency (see Figure 5(b) and Appendix G).

Evaluation metrics: We assess method effectiveness along three principal dimensions. **End-to-end retrieval latency**: Total query processing time. **Recall**: Result quality, defined as $\text{Recall} = \frac{|R_1 \cap R_2|}{|R_1|}$, where R_1 denotes the ground-truth result set and R_2 the retrieved set. **Uniformed-Cost**: To enable hardware-agnostic comparison, we adopt the rental-cost normalization from [113]. Results are reported as estimated cost per one million queries (\$/1M queries), using the hourly hardware rental rates in Fig. 5(b); our full cost model is detailed in Appendix G.

6.2 Effectiveness Evaluation

In this section, we evaluate the effectiveness of multiple LIKE retrieval methods across four datasets under four different workload

patterns. We assess the effectiveness and quality of these methods from two perspectives: efficiency and quality.

Efficiency Evaluation: We conducted a comprehensive efficiency evaluation on four diverse datasets (IMDB, TPC-H, WIKI, and Reddit), assessing performance through two key metrics: raw retrieval latency and a normalized query processing cost. As illustrated in Figure 6(a), SMILE consistently demonstrates superior retrieval latency. It achieves a $1.88 - 41.56\times$ speedup over trigram-based GIN indexes and surpasses both prefix/suffix B+-trees and full-table scans by two and three orders of magnitude, respectively. This performance stems from SMILE's design, which synergizes $O(\ell)$ -time decoding with a lightweight, hash-based candidate validation mechanism. This approach effectively circumvents the inherent bottlenecks of traditional methods, such as the $O(\log N)$ complexity of tree-based indexes or the performance degradation of GIN when wildcards disrupt trigram extraction. To ensure a fair comparison, particularly given SMILE's use of GPU hardware, we also introduce a normalized cost metric (USD per million queries) that accounts for hardware and energy expenditure. This provides an equitable assessment against CPU-based baselines, with results shown in Figure 6(b). On large-scale, complex datasets like IMDB and TPC-H, SMILE's efficiency translates into a substantial cost advantage, proving to be one to two orders of magnitude more cost-effective than all baselines. For instance, its cost on IMDB is significantly lower than that of the GIN index. However, on smaller datasets like Reddit, where the overhead of our model is less amortized, the cost advantage narrows. Under such smaller dataset, the highly-optimized GIN index is $2\times$ more cost effective than SMILE. In summary, while GIN remains a viable option for smaller workloads, SMILE's architecture delivers both superior speed and unparalleled cost-efficiency for the large-scale, complex pattern matching tasks that are the primary target of modern data systems.

Quality evaluation: Figure 7(a) evaluates LIKE search quality across four datasets and query patterns. While conventional deterministic indexes ensure perfect recall at a high computational cost, approximate methods like GIN(Apx) falter on most datasets. LLM-based approaches show performance scaling with model size; large models like DeepSeek-V3 (671B) achieve certain recall up to 0.8 on simple patterns (W1, W2), but their performance collapses (<0.1) on complex ones (W3, W4), revealing an inability to grasp data semantics in a database specific domain. In stark contrast, our compact SMILE model consistently delivers high recall (>0.9) across all workloads. This demonstrates the efficacy of a specialized, data-aware model for LIKE predicate evaluation, positioning SMILE as a robust solution where both accuracy and efficiency are critical.

End-to-End Performance on Realistic Workloads. Figure 7(b) presents a comprehensive end-to-end (E2E) evaluation of SMILE on complex queries from the TPC-H (T) and IMDB-JOB (J) benchmarks, selected for their challenging LIKE predicates. As shown in Figure 7(b)(i), SMILE consistently outperforms all baselines, achieving speedups of $1.18\times$ to $9.51\times$ over the next-best method. This demonstrates its practical utility in accelerating full query plans that include complex operations like joins and aggregations. Figure 7(b)(ii) confirms that this performance gain is achieved while maintaining a consistently high recall of over 95%. This evaluation also highlights the structural fragility of trigram-based approximate indexes. Notably, GIN(Apx) exhibits highly inconsistent




Figure 6: Evaluation on retrieval efficiency (Y-axis in log scale).




Figure 7: Evaluation on retrieval quality and end-to-end (E2E) performance.

recall—near-perfect on TPC-H queries but collapsing to near-zero on several IMDB-JOB templates. This discrepancy stems from the

query structure: TPC-H queries often terminate with a GROUP BY operation, which can mask row-level filtering errors in the final

1277 aggregated output. Conversely, the IMDB-JOB queries are direct
 1278 selections where any filtering inaccuracy immediately impacts the
 1279 final result set. SMILE's stable, high recall across both benchmarks
 1280 underscores its robustness. While the substantial speedup from
 1281 accelerating the LIKE predicate is naturally diluted by other sys-
 1282 tem costs (e.g., I/O, joins) in E2E queries, the fundamental operator
 1283 acceleration provides significant practical value. A more granular
 1284 analysis is available in Appendix H.

1285
 1286 **Takeaway:** SMILE delivers both high efficiency and high recall for
 1287 LIKE retrieval across diverse workloads. *Efficiency-wise*, it outper-
 1288 forms trigram indexes, B⁺-trees, and full scans, thanks to $O(\ell)$ decod-
 1289 ing and lightweight hash-based validation—avoiding the bottlenecks
 1290 of tree traversal, index degeneration, and LLM overhead. *Quality-*
 1291 *wise*, SMILE achieves over 0.9 recall on all patterns, far surpassing
 1292 LLMs while matching deterministic indexes. Its compact design en-
 1293 ables sub-20 ms inference without sacrificing accuracy, making it
 1294 practical for accelerating complex queries.

6.3 Scalability Evaluation


In this section, we empirically investigate the scalability of our approach through experiments, analyzing storage overhead, training cost, and inference cost.

Space scalability: Figure 8a evaluates the storage overhead of SMILE against state-of-the-art LLMs and traditional LIKE-optimized indexes. SMILE requires only $\approx 16\text{MB}$ of parameter storage, which is 3–5 orders of magnitude less than large LLMs (e.g., DeepSeek-V3, Qwen-32B). Compared to GIN, GIST, and B+-trees, SMILE reduces space consumption by 23 \times , 82 \times , and 28 \times in average. This efficiency arises from two factors: (1) our model's use of attention mechanisms to compactly encode workload distribution, and (2) its specialized design for database predicates, avoiding the domain-agnostic parameter bloat of general-purpose LLMs. These results confirm that SMILE achieves exceptional space efficiency without compromising task-specific effectiveness.

Training scalability: Fig. 8b evaluates SMILE's training-phase scalability against traditional indexes. On larger datasets (IMDB, TPC-H), SMILE's construction time is comparable to GIN and five times faster than GIST, though B+-tree remains fastest due to its batched, sorted-data indexing. GIN and GIST incur overhead from tokenizing column data into triplets and building inverted/general search trees over expanded token sets. SMILE exploits GPU parallelism via its attention mechanism, enabling concurrent gradient computation and updates for 10^4 tuple pairs per batch. Its compact parameter count further reduces convergence time. Consequently, SMILE achieves superior training scalability through hardware parallelism and attention efficiency.

Inference scalability: Fig. 9a shows end-to-end latency under varying sample budgets. Latency scales linearly with sample count: below 1ms for 8 samples, and close to 100ms for 512. At maximum sampling (512), this 0.1s latency is competitive with existing baselines under equivalent workloads, demonstrating superior throughput for high-volume inference.


Scalability with Cardinality. Figure 9b evaluates SMILE's de-
 1331 coding scalability with respect to query cardinality on the TPC-H
 1332 lineitem table. A fixed sample budget often suffices to achieve



(a) Scalability on space. (b) Scalability on building time.
 Figure 8: Scalability on space and building time(Y-axis in log scale).

high recall ($\geq 90\%$), confirming SMILE's efficiency. Moreover, the sample size needed to maintain high recall scales approximately linearly with result cardinality. This aligns with Theorem 5.2: higher-cardinality queries typically exhibit greater entropy and complexity, and the theorem predicts that the required sample size must scale accordingly—empirically observed here as a linear dependence on result size, thereby validating our theoretical bound.

Mitigating High-Entropy Challenges. High-entropy queries are inherently difficult, but our theory (Eq. 4) suggests a remedy: increasing model capacity, $|\theta|$. To validate this, we study the interplay among model size, query entropy, and recall. Figure 9c shows a representative result (Appendix C provide details). Entropy is calculated using Eq. 2. Two insights emerge: (1) for a fixed model size, recall degrades as entropy increases, confirming entropy as a performance boundary, as predicted; and (2) this boundary is malleable—increasing model capacity significantly boosts recall, especially in high-entropy regimes. This reveals a practical scaling law: the performance penalty from high entropy can be effectively offset by larger models, offering a principled trade-off between capacity and robustness on complex workloads. See Appendix C for a comprehensive analysis.



(a) Sample budget(Y-axis in log scale). (b) Query cardinality.

(c) Impact of model capacity across different query entropy regimes.

Figure 9: Scalability on budget, cardinality, entropy.



(a) SMILE (16MB) on TPC-H. (b) GPT-2 (124M) on NewsRoom.

Scalability with String Length. Figure 10 evaluates the scalability of our approach with respect to the length of the LIKE predicate. Figure 10a shows the performance of our standard 16MB SMILE model on the IMDB dataset for short predicates (up to 45 characters). As expected, latency scales linearly with predicate length, while recall remains consistently high. Figure 10b demonstrates our system's ability to handle extremely long predicates, a significant challenge for traditional DBMS [93, 110]. This experiment,

conducted on W4 of the News Room dataset, utilizes an extension of our prefix generation technique (Section 5.2) with a fine-tuned 124M GPT-2 model. As detailed in Appendix B.2, we find that a remarkably small number of generated prefix samples is sufficient for high accuracy. With just 8 samples, we achieve recall above 0.95 for predicates up to 8192 characters. This result is achieved with modest overhead: latency remains under 100ms, showcasing a highly favorable trade-off between accuracy and performance. A comprehensive analysis of the interplay between sample count, latency, and recall is provided in Appendix B.2.



(a) Generalization from 1% data. (b) Varying data proportion(W4).
Figure 11: Scalability on the billion-scale RedPajama dataset.


Scalability Across Billion-Scale Data. To evaluate the scalability and cost-effectiveness of our approach, we conducted experiments on the billion-record RedPajama dataset. First, we trained a 16MB SMILE model on 1% of the data and tested its generalization on an unseen, disjoint subset. As shown in Figure 11a, the model effectively learned the data distribution for simpler workloads, achieving over 0.9 recall for W2 with an inference sample size of 512. However, performance on the more complex workload W4 was lower, indicating the need for more training data. To determine the sufficient data scale, we trained models on progressively larger fractions of the dataset, focusing on the challenging W4 workload. Figure 11b shows that recall improves significantly as the training data proportion increases from 0.1% to 10%. Critically, we found that training on 10% of the data is sufficient to achieve excellent performance, yielding 0.9 recall with 512 samples. These results demonstrate that our method is both scalable and data-efficient. This finding corroborates that high performance does not require the full dataset or oversized models. Since query workload distributions are similar, the model can learn a generalizable mapping from a representative subset, making performance largely independent of absolute data scale once sufficiency is reached.

Takeaway: Our evaluation confirms SMILE is a practical and resource-efficient solution for LIKE queries, demonstrating superior scalability:
1. On Space: 3–5 orders of magnitude smaller than LLMs. **2. On Time:** Achieves competitive build times and comparable inference latency to traditional methods, while scaling gracefully to long predicates. **3. On Data Efficiency:** Generalizes on billion-scale data by training on only a small fraction (e.g., 10%), validating its efficiency.

6.4 Robustness Evaluation

The effectiveness of SMILE primarily derives from two key factors: the learned joint character distribution across database columns and the query patterns extracted from historical workloads. To systematically evaluate the robustness of SMILE under deviations from normal conditions, we conduct controlled experiments simulating both query drift and data drift scenarios using the realworld IMDB dataset (15M tuples with maximum string length attributes). This dataset presents significant robustness challenges due to its complex string patterns and scale.

Robustness to Data Drifts. We evaluate SMILE’s robustness to data drifts by simulating periodic data ingestion. The IMDB-name table is partitioned into five segments, and we test three variants: *Stale* (no updates), *FastUpd* (fine-tuned on 10 batches from each new partition), and *LongUpd* (fine-tuned on 100 batches).



(a) Update Scalability. (b) Recall Robustness.
Figure 12: Performance of SMILE under data drifts, showing the trade-off between update cost and recall.

As shown in Fig. 12, the results highlight an effective trade-off between performance and update cost. Remarkably, the *Stale* model maintains 84% recall on the final partition, demonstrating strong generalization from its initial training. The *FastUpd* strategy offers an efficient compromise, restoring recall close to 95% with a lightweight fine-tuning process that totals only 75 seconds. For maximum precision, *LongUpd* achieves near-perfect recall at a higher computational cost. This experiment confirms three key properties of SMILE: **(P1) Inherent Robustness:** It naturally generalizes to data with similar distributions, even without updates. **(P2) Efficient Adaptation:** Lightweight fine-tuning provides a cost-effective method to maintain high accuracy on evolving data. **(P3) Flexible Precision:** The architecture allows for more intensive retraining to achieve maximum precision when resources permit. These validate SMILE’s suitability for dynamic environments.





Figure 13: Robustness on wildcard patterns against inverted index.


Robustness on patterns. To evaluate robustness, we compare our method against inverted index across three workloads shown in Figure 13. The workloads are designed around keyword selectivity, where a keyword is defined as high-selectivity if its data frequency is below 0.5% and low-selectivity if above 2%. **Type-A** queries ("%"low%"") conjoin two common keywords. **Type-B** ("%"high%"") contains at least one rare keyword. **Type-C** ("%"word%"") contains single keyword with equal mix of high and low selectivity keywords. To simulate real-world complexity, we inject a varying number of "_" wildcards (0-7) into each pattern. The results highlight a clear performance trade-off. For Type-A queries, our generative approach is substantially faster, as it holistically processes the pattern, bypassing the GIN index’s expensive intersection of large posting lists. On Type-B and Type-C queries, performance is comparable since GIN can leverage a highly selective keyword or perform a direct lookup. Critically, Figure 13 shows that GIN’s latency degrades significantly with more "_" wildcards, as they fragment the trigrams the index relies on. In contrast, our method’s performance remains stable, demonstrating superior robustness to the complex and fragmented patterns that challenge traditional n-gram indexes. See Appendix D for a more analysis.

Robustness to Query Drifts. In Fig. 14, we investigate SMILE's OOD generalization by cross-evaluating on workloads with varying structural characteristics (W1-W4). Our analysis reveals two primary findings. First, a model trained on the diverse patterns of the W3 workload acts as a robust "generalist," adapting effectively to the realistic query structures of W1 and W2. This indicates that SMILE learns fundamental wildcard logic, not just template-specific features. However, generalization proves challenging when faced with the highly arbitrary and hard patterns of the W4 workload, which represents a significant structural shift and defines the model's performance boundaries. Second, and more importantly, this generalization gap can be substantially mitigated by increasing the inference-time sampling budget. As shown by comparing Fig. 14a and Fig. 14b, a larger sampling scale boosts recall by 10-20% on OOD templates. This highlights a key practical advantage of our approach: the ability to trade computational resources for enhanced robustness against unforeseen query drifts, effectively adapting to new patterns without retraining [35, 116, 120].



(a) Query OOD eval (S-32). (b) Query OOD eval (S-256).

Figure 14: Query drift evaluation.



(a) Decoding robustness. (b) Classification quality.

Figure 15: Decoding/Classification robustness.

Robustness of Different Decoding Methods. Fig. 15a evaluates the robustness of four decoding strategies for SMILE trained on templates from W1-W4 and tested on a disjoint set from the same wildcard workload distribution. Greedy decoding demonstrates poor generalization, frequently generating repetitive outputs and failing. Maximum likelihood and low-temperature ($T = 0.5$) decoding provide only slight improvements and remain ineffective on the more demanding W3 and W4 workloads. Conversely, high-temperature ($T = 1.4$) decoding achieves an optimal trade-off between quality and diversity, maintaining high recall across all workload patterns.

Robustness of the cardinality classifier: Fig. 15b assesses the classifier's ability to identify high-cardinality queries (threshold: 16) under varying wildcard patterns. Four classifier variants, each tuned to a specific workload pattern, achieve F1-scores >90% across all configurations. This performance arises from the SLM encoder's semantic representation, which maps query patterns into a latent space where high-cardinality queries become linearly separable. The resulting features enable the neural classifier to reliably identify and route such queries, enhancing processing efficiency.

Takeaway: SMILE exhibits strong robustness to both data and query drift. It generalizes effectively to evolving data, enabling a flexible trade-off between accuracy and fine-tuning cost. Unlike traditional inverted indexes, its performance remains stable across complex wildcard patterns. Meanwhile, SMILE mitigates query drift by trading

inference-time computation for generalization: increasing the sampling budget significantly improves zero-shot recall on unseen query templates, offering a practical alternative to expensive retraining.

7 RELATED WORK

Learned cardinality estimation on LIKE predicates: Cardinality estimation refers to the problem of predicting the result size of a query without executing it [57], leveraging either data distributions [87, 138, 139] or query patterns [67, 86]. Recently, learned approaches have become prominent in this domain, offering superior accuracy compared to traditional histogram or sampling-based methods [139]. In terms of string LIKE predicates, many works target estimating the cardinality of such predicates using learned models [14, 63, 75, 80]. However, there exists a fundamental difference between the cardinality estimation problem studied in prior work and our focus on accelerating LIKE predicates. While cardinality estimation aims to predict the number of matching tuples, LIKE predicate acceleration focuses on optimizing the end-to-end execution efficiency of queries involving LIKE operations. Prior methods addressing LIKE cardinality estimation do not address the core challenge of reducing the computational overhead inherent in pattern-matching operations, which is the central objective of our proposed solution.

Learned query acceleration: Recent research efforts increasingly replace traditional data structures with machine learning components [38, 51, 70, 71, 109, 126, 130]. Learned indexes revolutionize query processing by substituting logarithmic-time tree traversals with constant-time ML model inference [51, 70, 130]. For relational data, the Recursive Model Index (RMI) [70] employs multi-layer perceptrons to predict record locations, while the Fitting-Tree [51] achieves error-bounded performance through piecewise linear regression. Current learned string indexing approaches [117, 128, 137] optimize for variable-length storage and full/prefix matches, but they fail to support arbitrary LIKE wildcard patterns (e.g., %keyword%). This limitation arises from the complex challenge of modeling substring distributions under variable wildcard configurations. Our solution bridges this gap by adapting neural sequence-to-sequence models from NLP [16, 64, 118, 127]. Rather than relying on recursive tree structures, our approach resembles neural machine translation, transforming complex LIKE wildcard templates into candidates that are easy to verify. This also fundamentally differs from generative retrieval in NLP [54], which generates hypothetical documents for high-quality dense retrieval; in contrast, our work performs constrained database-specific generation to accelerate matching of complex LIKE patterns.

Efficient LIKE Processing & Approximate String Matching: The LIKE predicate is ubiquitous in databases [14, 65, 75] but hard to optimize, especially for wildcard-surrounded patterns (e.g., "%keyword%") [39, 124]. While B⁺-trees support prefix/suffix matches [33, 48], n-gram-based indexes (e.g., trigrams) are commonly used for general patterns via inverted [65, 101, 147] or tree structures [39]. Approximate substring matching, often based on edit distance [18, 25, 55, 82, 84, 136], addresses vague queries in interactive systems [17, 27, 62, 81]. Unlike prior work focusing on fixed-threshold, wildcard-free patterns [82], we accelerate LIKE execution by adopting Approximate Query Processing principles [24, 100], trading

1567
1568
1569
1570
1571
1572
1573
1574
1575
1576
1577
1578
1579
1580
1581
1582
1583
1584
1585
1586
1587
1588
1589
1590
1591
1592
1593
1594
1595
1596
1597
1598
1599
1600
1601
1602
1603
1604
1605
1606
1607
1608
1609
1610
1611
1612
1613
1614
1615
1616
1617
1618
1619
1620
1621
1622
1623
1624

completeness for speed. Our method returns a high-quality subset of guaranteed-correct matches, avoiding costly full scans on string columns when prefix-index optimization fails, especially given LIKE's prevalence in benchmarks like TPC-H (Q2,9,13,14,20) [125].

8 CONCLUSION

This paper introduces SMILE, a novel framework that leverages a small language model to accelerate complex LIKE queries with complex wildcards. Our evaluation shows that SMILE outperforms traditional indexing structures by up to three orders of magnitude in execution time and remains robust to shifts in both data and query distributions, highlighting its potential for modern DBMS. Our work suggests several promising directions for future research. First, the core idea: casting complex query execution as neuron translation—can be extended beyond wildcard matching to other costly operations, such as regular expression evaluation and similarity joins. Second, while our current transformer-based SLM is effective against LLMs, more lightweight or database-native architectures may reduce inference overhead, particularly for high-entropy workloads where existing decoding struggles. In the future, we aim to develop more effective learned structures, enabling controllable trade-offs between performance and accuracy across diverse situations.

REFERENCES

- [1] [n. d.]. dataset. <https://bulk-text-to-sql-benchmark.github.io/>
- [2] [n. d.]. dataset. <https://huggingface.co/datasets/zhanghanchong/css/tree/main>
- [3] [n. d.]. dataset. <https://github.com/gregrahn/join-order-benchmark>
- [4] [n. d.]. dataset. <https://github.com/Ffunkytao/LogicCat>
- [5] [n. d.]. dataset. <https://github.com/ygan/SpiderSS-SpiderCG/tree/main>
- [6] [n. d.]. dataset. <https://github.com/peterbaile/beaver>
- [7] 2023. Information technology – Database languages – SQL – Part 2: Foundation (SQL/Foundation). <https://webstore.iec.ch/en/publication/86046>
- [8] 2024. SMILE: Realworld_84047LIKE.csv. Anonymous 4open.science repository. https://anonymous.4open.science/r/SMILE-0D8B/Realworld_84047LIKE.csv Accessed: 2025-10-31.
- [9] 2025. morfessor. <https://morfessor.readthedocs.io/en/latest/> Accessed: 2025-10-31.
- [10] 2025. Pyphen. <https://pyphen.org/> Accessed: 2025-10-31.
- [11] Vanna AI. 2023. Vanna: LLM-powered SQL Assistant (Examples/Datasets). GitHub repository. <https://github.com/vanna-ai/vanna>
- [12] Gene M. Amdahl. 1967. Validity of the Single-Processor Approach to Achieving Large Scale Computing Capabilities. In Proceedings of the AFIPS Spring Joint Computer Conference, Vol. 30. AFIPS Press, 483–485.
- [13] AutoDL. 2024. AutoDL GPU Marketplace. <https://www.autodl.com/market/list> Accessed: 2024-06-15.
- [14] Mehmet Aytimir, Silvan Reiner, Leonard Wörteler, Theodoros Chondrogiannis, and Michael Grossniklaus. 2024. LPLM: A Neural Language Model for Cardinality Estimation of LIKE-Queries. Proc. ACM Manag. Data 2, 1, Article 54 (March 2024), 25 pages. <https://doi.org/10.1145/3639309>
- [15] Mehmet Aytimir, Silvan Reiner, Leonard Wörteler, Theodoros Chondrogiannis, and Michael Grossniklaus. 2024. LPLM: A Neural Language Model for Cardinality Estimation of LIKE-Queries. Proc. ACM Manag. Data 2, 1 (2024), 1–25.
- [16] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. 2014. Neural Machine Translation by Jointly Learning to Align and Translate. CoRR abs/1409.0473 (2014). <https://api.semanticscholar.org/CorpusID:11212020>
- [17] Holger Bast, Debapriyo Majumdar, Ralf Schenkel, Martin Theobald, and Gerhard Weikum. 2006. IO-Top-k: index-access optimized top-k query processing. In Proceedings of the 32nd International Conference on Very Large Data Bases (Seoul, Korea) (VLDB '06). VLDB Endowment, 475–486.
- [18] Alexander Behm, Shengyue Ji, Chen Li, and Jiaheng Lu. 2009. Space-Constrained Gram-Based Indexing for Efficient Approximate String Search. In Proceedings of the 2009 IEEE International Conference on Data Engineering (ICDE '09). IEEE Computer Society, USA, 604–615. <https://doi.org/10.1109/ICDE.2009.32>
- [19] BIRD Team. 2023. BIRD: Text-to-Efficient-SQL Benchmark. Project website. <https://bird-bench.github.io/> Focus on both correctness and efficiency of generated SQL.
- [20] Rishi Bommasani, Drew A. Hudson, Ehsan Adeli, Russ Altman, Simran Arora, Sydney von Arx, Michael S. Bernstein, Jeannette Bohg, Antoine Bosselut, Emma Brunskill, Erik Brynjolfsson, S. Buch, Dallas Card, Rodrigo Castellon, Niladri S. Chatterji, Annie S. Chen, Kathleen A. Creel, Jared Davis, Dora Demszky, Chris Donahue, Moussa Doumbouya, Esin Durmus, Stefano Ermon, John Etchemendy, Chelsea Finn, Trevor Gale, Lauren E.Gillespie, Karan Goel, Noah D. Goodman, Shelby Grossman, Neel Guha, Tatsunori Hashimoto, Peter Henderson, John Hewitt, Daniel E. Ho, Jenny Hong, Kyle Hsu, Jing Huang, Thomas F. Icard, Saahil Jain, Dan Jurafsky, Pratyusha Kalluri, Siddharth Karamcheti, Geoff Keeling, Fereshte Khani, O. Khattab, Pang Wei Koh, Mark S.Krass, Ranjay Krishna, Rohith Kuditipudi, Ananya Kumar, Faisal Ladha, Mina Lee, Tony Lee, Jure Leskovec, Isabelle Levent, Xiang Lisa Li, Xuechen Li, Tengyu Ma, Ali Malik, Christopher D. Manning, Suvir P. Mirchandani, Eric Mitchell, Zanele Munyikwa, Suraj Nair, Avaniika Narayan, Deepak Narayanan, Benjamin Newman, Allen Nie, Juan Carlos Niebles, Hamed Nilforoshan, J. F.Nyarko, Giray Ogut, Laurel Orr, Isabel Papadimitriou, Joon Sung Park, Chris Piech, Eva Portelance, Christopher Potts, Aditi Raghunathan, Robert Reich, Hongyu Ren, Frieda Rong, Yusuf H. Roohani, Camilo Ruiz, Jack Ryan, Christopher R'e, Dorsa Sadigh, Shiori Sagawa, Keshav Santhanam, Andy Shih, Krishna Parasuram Srinivasan, Alex Tamkin, Rohan Taori, Armin W. Thomas, Florian Tram'r, Rose E.Wang, William Wang, Bohan Wu, Jiajun Wu, Yuhuai Wu, Sang Michael Xie, Michihiro Yasunaga, Jiaxuan You, Matei A. Zaharia, Michael Zhang, Tianyi Zhang, Xikun Zhang, Yuhui Zhang, Lucia Zheng, Kaityn Zhou, and Percy Liang. 2021. On the Opportunities and Risks of Foundation Models. *ArXiv* (2021). <https://crfm.stanford.edu/assets/report.pdf>
- [21] Felix Brauer, Ralf Rieger, Adrian Mocan, Craig A. Knoblock, and Pedro Szekely. 2011. Enabling Information Extraction by Inference of Regular Expressions from Sample Entities. In CIKM. 1285–1290.
- [22] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language models are few-shot learners. *arXiv preprint arXiv:2005.14165* (2020).
- [23] Chen Chai, Lei Cao, Guoliang Li, Samuel Madden, and Nan Tang. 2020. Human-in-the-loop Outlier Detection. In SIGMOD. 19–33.
- [24] Kaushik Chakrabarti, Minos Garofalakis, Rajeev Rastogi, and Kyuseok Shim. 2000. Approximate Query Processing Using Wavelets. In Proceedings of the 26th International Conference on Very Large Data Bases (VLDB). Morgan Kaufmann, 111–122.
- [25] S. Chaudhuri, K. Ganjam, V. Ganti, and R. Motwani. 2003. Robust and Efficient Fuzzy Match for Online Data Cleaning. In Proceedings of the 2003 ACM SIGMOD International Conference on Management of Data (SIGMOD). 313–324.
- [26] Surajit Chaudhuri and Luis Gravano. 1999. Evaluating Top-k Selection Queries. In Proceedings of the 25th International Conference on Very Large Data Bases (VLDB '99). Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 397–410.
- [27] S. Chaudhuri and R. Kaushik. 2009. Extending Autocompletion to Tolerate Errors. In Proceedings of the 2009 ACM SIGMOD International Conference on Management of Data (SIGMOD '09). ACM. <https://doi.org/10.1145/1559845.1559868>
- [28] Hui Chen, Xiaoyan Liu, Dingxiao Yin, and Jie Tang. 2017. A survey on dialogue systems: Recent advances and new frontiers. *Acm Sigkdd Explorations Newsletter* 19, 2 (2017), 25–35.
- [29] Qian Chen, Arka Banerjee, Çağatay Demirpal, Samuel Madden, and Nesime Tatbul. 2023. Data Extraction via Semantic Regular Expression Synthesis. *Proc. ACM Program. Lang.* 7, OOPSLA2 (2023), 1848–1877.
- [30] Cherry Servers. 2024. AMD EPYC 9654 Dedicated Server Pricing. <https://www.cherryservers.com/pricing/dedicated-servers/amd-epyc-9654?region=de-frankfurt&billing=37> Accessed: 2024-06-15.
- [31] Transaction Processing Performance Council and Electrum. 2018. TPC-H DBGEN Data Generator. GitHub repository. <https://github.com/electrum/tpch-dbgem> TPC-H specification data generator for analytic SQL benchmarking.
- [32] Michele Dallachiesa, Ahmed Ebaid, Ahmed Eldawy, Ahmed Elmagarmid, Mourad Ouazzani, Dimitris Papadias, and Nan Tang. 2013. NADEF: A Commodity Data Cleaning System. In SIGMOD. 541–552.
- [33] René De La Briandais. 1959. File searching using variable length keys. In Papers Presented at the the March 3–5, 1959, Western Joint Computer Conference (San Francisco, California) (IRE-AIEE-ACM '59 (Western)). Association for Computing Machinery, New York, NY, USA, 295–298. <https://doi.org/10.1145/1457838.1457895>
- [34] DeepSeek-AI. 2024. DeepSeek-V2: A Strong, Economical, and Efficient Mixture-of-Experts Language Model. *arXiv:2405.04434 [cs.CL]*
- [35] DeepSeek-AI. 2025. DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning. *arXiv:2501.12948 [cs.CL]* <https://arxiv.org/abs/2501.12948>
- [36] Dong Deng, Guoliang Li, and Jianhua Feng. 2012. An Efficient Trie-based Method for Approximate Entity Extraction with Edit-Distance Constraints. In Proceedings of the 2012 IEEE 28th International Conference on Data

- 1741 Engineering (ICDE '12). IEEE Computer Society, USA, 762–773. <https://doi.org/10.1109/ICDE.2012.29>
- 1742 [37] Jacob Devlin, Michael-W Chang, Kenton Lee, and Kristina Toutanova. 2018. BERT: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805* (2018).
- 1743 [38] Jialin Ding, Ryan Marcus, Andreas Kipf, Vikram Nathan, Aniruddha Rusimha, Kapil Vaitya, Alexander van Renen, and Tim Kraska. 2022. SageDB: An Instance-Optimized Data Analytics System. *Proc. VLDB Endow.* 15, 13 (Sept. 2022), 4062–4078. <https://doi.org/10.14778/3565838.3565857>
- 1744 [39] PostgreSQL Documentation. [n. d.]. F.33. pg_trgm — support for similarity of text using trigram matching. <https://www.postgresql.org/docs/current/pgtrgm.html>. Accessed: 2025-03-07.
- 1745 [40] Qingxiu Dong, Lei Li, Damai Dai, Ce Zheng, Jingyuan Ma, Rui Li, Heming Xia, Jingjing Xu, Zhiyong Wu, Baobao Chang, Xu Sun, Lei Li, and Zhifang Sui. 2024. A Survey on In-context Learning. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen (Eds.). Association for Computational Linguistics, Miami, Florida, USA, 1107–1128. <https://doi.org/10.18653/v1/2024.emnlp-main.64>
- 1746 [41] Eduard Dragut, Fang Fang, Prasad Sistla, Clement Yu, and Weiyi Meng. 2009. Stop word and related problems in web interface integration. *Proc. VLDB Endow.* 2, 1 (Aug. 2009), 349–360. <https://doi.org/10.14778/1687627.1687667>
- 1747 [42] Lynne G Duncan. 2018. Language and Reading: the Role of Morpheme and Phoneme Awareness. *Current Developmental Disorders Reports* 5, 4 (2018), 226–234. <https://doi.org/10.1007/s40474-018-0153-2>
- 1748 [43] DBA Stack Exchange. [n. d.]. LIKE Query Optimization. <https://dba.stackexchange.com/questions/203748/like-query-optimization>. Accessed: 2025-03-07.
- 1749 [44] Ju Fan, Zihui Gu, Songyue Zhang, Yuxin Zhang, Zui Chen, Lei Cao, Guoliang Li, Samuel Madden, Xiaoyong Du, and Nan Tang. 2024. Combining small language models and large language models for zero-shot NL2SQL. *Proceedings of the VLDB Endowment* 17, 11 (2024), 2750–2763.
- 1750 [45] Ju Fan and Guoliang Li. 2018. Human-in-the-loop Rule Learning for Data Integration. *IEEE Data Eng. Bull.* 41, 2 (2018), 104–115.
- 1751 [46] Wenfei Fan, Jianzhong Li, Shuai Ma, Nan Tang, and Wenyuan Yu. 2011. Interaction between record matching and data repairing. (2011), 469–480. <https://doi.org/10.1145/1989323.1989373>
- 1752 [47] Wanxiang Feng, Dong Han, Shuai Zhang, Linzhang Wang, Yu Wang, Minghui Qiu, Di Wang, Hui Wang, and Juanzi Li. 2020. CodeBERT: A Pre-Trained Model for Programming and Natural Language Processing. *arXiv preprint arXiv:2002.08155* (2020). <https://arxiv.org/abs/2002.08155>
- 1753 [48] Paolo Ferragina and Roberto Grossi. 1999. The string B-tree: a new data structure for string search in external memory and its applications. *J. ACM* 46, 2 (March 1999), 236–280. <https://doi.org/10.1145/301970.301973>
- 1754 [49] Robert M. French. 1999. Catastrophic Forgetting in Connectionist Networks. *Trends in Cognitive Sciences* 3, 4 (1999), 128–135.
- 1755 [50] Victoria Fromkin, Robert Rodman, and Nina Hyams. 2013. *An Introduction to Language* (10th ed.). Wadsworth Cengage Learning, Boston. 25–55 pages. PDF lecture extract available at <https://www3.uji.es/~ruiz/0903/Lects/Fromkin-etal13-Ch2-Morphology.pdf>.
- 1756 [51] Alex Galakatos, Michael Markovitch, Carsten Binnig, Rodrigo Fonseca, and Tim Kraska. 2019. Fiting-tree: A data-aware index structure. In *Proceedings of the 2019 International Conference on Management of Data*. 1189–1206.
- 1757 [52] Yuhui Gan and collaborators. 2021. Spider-Syn: Synthetic Paraphrases for Spider. GitHub repository. https://github.com/yghan/Spider-Syn/blob/main/Spider-Syn/train_spider.json
- 1758 [53] Yuhui Gan, Xinyun Zhang, Yujia Ou, Alane Suhr, Jiaqi Guo, Caiming Xiong, Yu Su, Tao Yu, Greg Durrett, Junyi Jessy Zhu, Ran Xu, Wenhan Wang, Nelson F. Wong, Victor Zhong, Dragomir Radev, et al. 2021. Exploring Underexplored Limitations of Cross-Domain Text-to-SQL Generalization. In *Proceedings of EMNLP 2021*. <https://aclanthology.org/2021.emnlp-main.702.pdf> Introduces Spider-DK dataset; official code/data: <https://github.com/yghan/Spider-DK>.
- 1759 [54] Luyu Gao, Xueguang Ma, Jimmy Lin, and Jamie Callan. 2023. Precise Zero-Shot Dense Retrieval without Relevance Labels. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 1762–1777. <https://doi.org/10.18653/v1/2023.acl-long.99>
- 1760 [55] L. Gravano, P. G. Ipeirotis, H. V. Jagadish, N. Koudas, S. Muthukrishnan, and D. Srivastava. 2001. Approximate String Joins in a Database (Almost) for Free. In *Proceedings of the 27th International Conference on Very Large Data Bases (VLDB)*, 491–500.
- 1761 [56] Max Grusky, Mor Naaman, and Yoav Artzi. 2018. NewsRoom: A Large-Scale News Summarization Dataset and Abstractive Summarization Baseline. <http://lil.nlp.cornell.edu/resources/newsroom/r8625bda324/newsroom-release.tar> Dataset release tarball, accessed 2025-10-24.
- 1762 [57] Yuxing Han, Ziniu Wu, Peizhi Wu, Rong Zhu, Jingyi Yang, Liang Wei Tan, Kai Zeng, Gao Cong, Yanzhao Qin, Andreas Pfandler, Zhengping Qian, Jingren
- 1763 Zhou, Jiangneng Li, and Bin Cui. 2021. Cardinality estimation in DBMS: a comprehensive benchmark evaluation. *Proc. VLDB Endow.* 15, 4 (dec 2021), 752–765. <https://doi.org/10.14778/3503585.3503586>
- 1764 [58] Yunzhong He, Yuxin Tian, Mengjiao Wang, Feier Chen, Licheng Yu, Mao-long Tang, Congcong Chen, Ning Zhang, Bin Kuang, and Arul Prakash. 2023. Que2Engage: Embedding-based Retrieval for Relevant and Engaging Products at Facebook Marketplace. In *Companion Proceedings of the ACM Web Conference 2023* (Austin, TX, USA) (*WWW '23 Companion*). Association for Computing Machinery, New York, NY, USA, 386–390. <https://doi.org/10.1145/3543873.3584633>
- 1765 [59] Johannes Hürluppi and collaborators. 2021. SEDE: Stack Exchange Data Explorer NL→SQL Dataset. GitHub repository. <https://github.com/hirupert/sede> Real-world NL-SQL pairs from Stack Exchange Data Explorer.
- 1766 [60] IBM Corporation 2023. *IBM Db2 13 for z/OS SQL Reference*. IBM Corporation. Sec. “LIKE predicate”, https://www.ibm.com/docs/en/SSEPEK_13.0.0/pdf/db2z_13_sqlrefbook.pdf.
- 1767 [61] IMDB. [n. d.]. Latest IMDB-Name(2025). <https://datasets.imdbws.com/>. Accessed: 2025-10-21.
- 1768 [62] S. Ji, G. Li, C. Li, and J. Feng. 2009. Efficient Interactive Fuzzy Keyword Search. In *Proceedings of the 18th International Conference on World Wide Web (WWW '09)*. ACM. <https://doi.org/10.1145/1526709.1526718>
- 1769 [63] Liang Jin and Chen Li. 2005. Selectivity estimation for fuzzy string predicates in large data sets. In *Proceedings of the 31st International Conference on Very Large Data Bases (Trondheim, Norway) (VLDB '05)*. VLDB Endowment, 397–408.
- 1770 [64] Melvin Johnson, Mike Schuster, Quoc V. Le, Maxim Krikun, Yonghui Wu, Zhifeng Chen, Nikhil Thorat, Fernanda Viégas, Martin Wattenberg, Greg Corrado, Macduff Hughes, and Jeffrey Dean. 2017. Google’s Multilingual Neural Machine Translation System: Enabling Zero-Shot Translation. *Transactions of the Association for Computational Linguistics* 5 (2017), 339–351. https://doi.org/10.1162/tacl_a_00065
- 1771 [65] Younghoon Kim, Hyoungmin Park, Kyuseok Shim, and Kyoung-Gu Woo. 2013. Efficient processing of substring match queries with inverted variable-length gram indexes. *Information Sciences* 244 (2013), 119–141. <https://doi.org/10.1016/j.ins.2013.04.037>
- 1772 [66] Younghoon Kim and Kyuseok Shim. 2013. Efficient top-k algorithms for approximate substring matching. In *Proceedings of the 2013 ACM SIGMOD International Conference on Management of Data (New York, New York, USA) (SIGMOD '13)*. Association for Computing Machinery, New York, NY, USA, 385–396. <https://doi.org/10.1145/2463676.2465324>
- 1773 [67] Andreas Kipf, Thomas Kipf, Bernhard Radke, Viktor Leis, Peter A. Boncz, and Alfons Kemper. 2018. Learned Cardinalities: Estimating Correlated Joins with Deep Learning. *ArXiv abs/1809.00677* (2018). <https://api.semanticscholar.org/CorpusID:52154172>
- 1774 [68] James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, Demis Hassabis, Claudia Clopath, Dharshan Kumaran, and Raia Hadsell. 2017. Overcoming Catastrophic Forgetting in Neural Networks. *Proceedings of the National Academy of Sciences* 114, 13 (2017), 3521–3526.
- 1775 [69] Christopher Kosten and collaborators. 2023. ScienceBenchmark Dataset. GitHub repository. https://github.com/cskosten/sciencebenchmark_dataset
- 1776 [70] Tim Kraska, Alex Beutel, Ed H. Chi, Jeffrey Dean, and Neoklis Polyzotis. 2018. The case for learned index structures. In *Proceedings of the 2018 international conference on management of data*. 489–504.
- 1777 [71] Ani Kristo, Kapil Vaitya, Ugur Çetintemel, Sanchit Misra, and Tim Kraska. 2020. The Case for a Learned Sorting Algorithm. In *Proceedings of the 2020 ACM SIGMOD International Conference on Management of Data (Portland, OR, USA) (SIGMOD '20)*. Association for Computing Machinery, New York, NY, USA, 1001–1016. <https://doi.org/10.1145/3318464.3389752>
- 1778 [72] Jonathan K. Kummerfeld and collaborators. 2018. Advising: Academic Advising NL→SQL Set. *text2sql-data* subset (GitHub). <https://github.com/jkkummerfeld/text2sql-data/blob/master/data/advising.json>
- 1779 [73] Jonathan K. Kummerfeld and collaborators. 2018. *text2sql-data: A Collection of Text-to-SQL Datasets*. GitHub repository. <https://github.com/jkkummerfeld/text2sql-data>
- 1780 [74] Toshitaka Kuwa, Shigehiko Schamoni, and Stefan Riezler. 2020. Embedding Meta-Textual Information for Improved Learning to Rank. In *Proceedings of the 28th International Conference on Computational Linguistics*, Donia Scott, Nuria Bel, and Chengqing Zong (Eds.). International Committee on Computational Linguistics, Barcelona, Spain (Online), 5558–5568. <https://doi.org/10.18653/v1/2020.coling-main.487>
- 1781 [75] Suyong Kwon, Kyuseok Shim, and Woohwan Jung. 2025. Cardinality Estimation of LIKE Predicate Queries using Deep Learning. *Proc. ACM Manag. Data* 3, 1, Article 20 (Feb. 2025), 26 pages. <https://doi.org/10.1145/3709670>
- 1782 [76] Suyong Kwon, Kyuseok Shim, and Woohwan Jung. 2025. Cardinality Estimation of LIKE Predicate Queries using Deep Learning. *Proc. ACM Manag. Data* 3, 1

- 1857 (2025), 1–26.
- 1858 [77] SUDA-LA Lab. 2021. SeSQL: Chinese Text-to-SQL Datasets. GitHub repository. <https://github.com/SUDA-LA/SeSQL/>
- 1859 [78] Chia-Hsuan Lee. 2020. KaggleDBQA Examples (Greater Manchester Crime). GitHub example file. <https://github.com/Chia-Hsuan-Lee/KaggleDBQA/blob/main/examples/GreaterManchesterCrime.json>
- 1860 [79] Gyubok Lee, Hyeonji Hwang, Seongsu Bae, Yeonsu Kwon, Woncheol Shin, Seongju Yang, Minjoon Seo, Jong-Yeup Kim, and Edward Choi. 2022. EHRSQL: A Practical Text-to-SQL Benchmark for Electronic Health Records. In *NeurIPS 2022 Datasets and Benchmarks Track*. <https://github.com/glee4810/EHRSQL> Benchmark for Text-to-SQL on EHR data.
- 1861 [80] Hongrae Lee, Raymond T. Ng, and Kyuseok Shim. 2009. Approximate substring selectivity estimation. In *Proceedings of the 12th International Conference on Extending Database Technology: Advances in Database Technology* (Saint Petersburg, Russia) (EDBT '09). Association for Computing Machinery, New York, NY, USA, 827–838. <https://doi.org/10.1145/1516360.1516455>
- 1862 [81] Chengkai Li, Kevin Chen-Chuan Chang, Ihab F. Ilyas, and Sumin Song. 2005. RankSQL: query algebra and optimization for relational top-k queries. In *Proceedings of the 2005 ACM SIGMOD International Conference on Management of Data* (Baltimore, Maryland) (SIGMOD '05). Association for Computing Machinery, New York, NY, USA, 131–142. <https://doi.org/10.1145/1066157.1066173>
- 1863 [82] Chen Li, Bin Wang, and Xiaochun Yang. 2007. VGRAM: improving performance of approximate queries on string collections using variable-length grams. In *Proceedings of the 33rd International Conference on Very Large Data Bases* (Vienna, Austria) (VLDB '07). VLDB Endowment, 303–314.
- 1864 [83] Guoliang Li. 2017. Human-in-the-loop Data Integration. *Proc. VLDB Endow.* 10, 12 (2017), 2006–2017.
- 1865 [84] Guoliang Li, Dong Deng, and Jianhua Feng. 2011. Faerie: efficient filtering algorithms for approximate dictionary-based entity extraction. In *Proceedings of the 2011 ACM SIGMOD International Conference on Management of Data* (Athens, Greece) (SIGMOD '11). Association for Computing Machinery, New York, NY, USA, 529–540. <https://doi.org/10.1145/1989323.1989379>
- 1866 [85] Guoliang Li, Shengyu Ji, Chen Li, and Jianhua Feng. 2011. Efficient fuzzy full-text type-ahead search. *The VLDB Journal* 20, 4 (August 2011), 617–640. <https://doi.org/10.1007/s00778-011-0218-x>
- 1867 [86] Pengfei Li, Wenqing Wei, Rong Zhu, Bolin Ding, Jingren Zhou, and Hua Lu. 2023. ALFCE: An Attention-based Learned Cardinality Estimator for SPJ Queries on Dynamic Workloads. *Proc. VLDB Endow.* 17, 2 (oct 2023), 197–210. <https://doi.org/10.14778/3626292.3626302>
- 1868 [87] Yingze Li, Hongzhi Wang, and Xianglong Liu. 2024. One Seed, Two Birds: A Unified Learned Structure for Exact and Approximate Counting. *Proc. ACM Manag. Data* 2, 1, Article 15 (mar 2024), 26 pages. <https://doi.org/10.1145/3639270>
- 1869 [88] Zeyu Li, Hongzhi Wang, Wei Shao, Jianzhong Li, and Hong Gao. 2016. Repairing Data through Regular Expressions. *Proc. VLDB Endow.* 9, 5 (2016), 432–443. <https://doi.org/10.14778/2876473.2876478>
- 1870 [89] Zeyu Li, Hongzhi Wang, Wei Shao, Jianzhong Li, and Hong Gao. 2016. Repairing Data through Regular Expressions. *Proc. VLDB Endow.* 9, 5 (2016), 432–443.
- 1871 [90] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. 2019. Roberta: A robustly optimized bert pretraining approach. *arXiv preprint arXiv:1907.11692* (2019).
- 1872 [91] Yurong Liu, Eduardo Pena, Aécio Santos, Eden Wu, and Juliana Freire. 2024. Magneto: Combining Small and Large Language Models for Schema Matching. *arXiv preprint arXiv:2412.08194* (2024). <https://arxiv.org/pdf/2412.08194.pdf>
- 1873 [92] Yev Meyer, Marjan Emadi, Dhruv Nathawani, Lipika Ramaswamy, Kendrick Boyd, Maarten Van Segbroeck, Matthew Grossman, Piotr Mlocek, and Drew Newberry. 2024. Synthetic-Text-To-SQL: A synthetic dataset for training language models to generate SQL queries from natural language prompts. <https://huggingface.co/datasets/gretelai/synthetic-text-to-sql>
- 1874 [93] Microsoft 2023. LIKE (Transact-SQL). Microsoft. <https://learn.microsoft.com/en-us/sql/t-sql/language-elements/like-transact-sql?view=sql-server-ver17>.
- 1875 [94] Qingkai Min, Yuefeng Shi, and Yue Zhang. 2019. A Pilot Study for Chinese SQL Semantic Parsing. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*. 3643–3649.
- 1876 [95] Sewon Min, Mike Lewis, Luke Zettlemoyer, and Hannaneh Hajishirzi. 2022. MetaCL: Learning to Learn In Context. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, Marine Carpuat, Marie-Catherine de Marneffe, and Ivan Vladimir Meza Ruiz (Eds.). Association for Computational Linguistics, Seattle, United States, 2791–2809. <https://doi.org/10.18653/v1/2022-naacl-main.201>
- 1877 [96] Colin Morris. 2017. Reddit Usernames. <https://www.kaggle.com/datasets/colinmorris/reddit-usernames>. Accessed: 2025-03-18.
- 1878 [97] John X. Morris, Chawin Sitawarin, Chuan Guo, Narine Kokhlikyan, G. Edward Suh, Alexander M. Rush, et al. 2025. How much do language models memorize?
- 1879 [98] Brent Ozar. [n. d.]. Sargability: Why %string% Is Slow. <https://www.brentozar.com/archive/2010/06/sargable-why-string-is-slow/>. Accessed: 2025-03-07.
- 1880 [99] Rohit Shantanam Patil. 2019. SQL-NLP Paired Dataset (NLP_Dataset.csv). GitHub repository. https://github.com/rohitshantanampatil/sql-nlp/blob/master/utils/NLP_Dataset.csv
- 1881 [100] Jialin Peng, Donghui Zhang, and Feifei Li. 2018. AQP++: Connecting Approximate Query Processing With Aggregate Precomputation for Interactive Analytics. In *Proceedings of the 2018 ACM SIGMOD International Conference on Management of Data*. ACM, 1093–1107.
- 1882 [101] PostgreSQL Core Team. 2025. GIN Indexes. <https://www.postgresql.org/docs/current/gin.html>. Accessed: 2025-03-18.
- 1883 [102] Abdullah Qahtan, Nan Tang, Mourad Ouzzani, and Michael Stonebraker. 2020. Pattern Functional Dependencies for Data Cleaning. *Proc. VLDB Endow.* 13, 5 (2020), 684–697.
- 1884 [103] Alec Radford, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. 2018. Improving language understanding by generative pre-training. (2018).
- 1885 [104] Alec Radford, Jeffrey Wu, Rewon Child, David Chen, Dario Amodei, and Ilya Sutskever. 2019. Language models are unsupervised multitask learners. *OpenAI blog* 1, 8 (2019).
- 1886 [105] Vijayshankar Raman and Joseph M. Hellerstein. 2001. Potter's Wheel: An Interactive Data Cleaning System. In *VLDB*. 381–390.
- 1887 [106] Baidu Research. 2020. DuSQL: A Large-Scale Chinese Text-to-SQL Dataset. AI Studio dataset. <https://aistudio.baidu.com/datasetdetail/102167/1>
- 1888 [107] David Rolnick, Arun Ahuja, Jonathan Schwarz, Timothy P. Lillicrap, and Gregory Wayne. 2019. Experience Replay for Continual Learning. *NeurIPS* 32 (2019).
- 1889 [108] Olli Räsänen, Reima Laaksonen, and Mikko Kurimo. 2018. Pre-linguistic segmentation of speech into syllable-like units. *Cognition* 171 (2018), 130–150. <https://doi.org/10.1016/j.cognition.2017.11.003>
- 1890 [109] Ibrahim Sabek, Kapil Vaidya, Dominik Horn, Andreas Kipf, Michael Mitzenmacher, and Tim Kraska. 2022. Can Learned Models Replace Hash Functions? *Proc. VLDB Endow.* 16, 3 (Nov. 2022), 532–545. <https://doi.org/10.14778/3570690.3570702>
- 1891 [110] SAP SE 2023. SAP IQ 16.0 SP04 Reference: Building Blocks. SAP SE. PDF p. 80, <https://infocenter.sybase.com/help/topic/com.sybase.infocenter.dc38151.1604/doc/pdf/qrefbb.pdf>
- 1892 [111] Gaurav Saxena, Mohammad Rahman, Naresh Chainani, Chunbin Lin, George Caragea, Fahim Chowdhury, Ryan Marcus, Tim Kraska, Ippokratios Pandis, and Balakrishnan (Murali) Narayanaswamy. 2023. Auto-WLM: Machine Learning Enhanced Workload Management in Amazon Redshift. In *Companion of the 2023 International Conference on Management of Data* (Seattle, WA, USA) (SIGMOD '23). Association for Computing Machinery, New York, NY, USA, 225–237. <https://doi.org/10.1145/3555041.3589677>
- 1893 [112] seekly. 2024. SynSQL-2.5M: Synthetic NL→SQL Pairs. ModelScope dataset. <https://www.modelscope.cn/datasets/seekly/SynSQL-2.5M>
- 1894 [113] Anil Shanbhag, Samuel Madden, and Xiangyao Yu. 2020. A Study of the Fundamental Performance Characteristics of GPUs and CPUs for Database Analytics. In *Proceedings of the 2020 ACM SIGMOD International Conference on Management of Data* (Portland, OR, USA) (SIGMOD '20). Association for Computing Machinery, New York, NY, USA, 1617–1632. <https://doi.org/10.1145/3318464.3380595>
- 1895 [114] Saurabh Shetty, Saravanan Thirumuruganathan, Nick Koudas, and Divesh Srivastava. 2020. Astrid: Accurate Selectivity Estimation for String Predicates Using Deep Learning. *Proc. VLDB Endow.* 14, 4 (2020), 684–697.
- 1896 [115] Koustuv Sinha, Robin Jia, Dieuwke Hupkes, Joelle Pineau, Adina Williams, and Douwe Kiela. 2021. Masked Language Modeling and the Distributional Hypothesis: Order Word Matters Pre-training for Little. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih (Eds.). Association for Computational Linguistics, Online and Punta Cana, Dominican Republic, 2888–2913. <https://doi.org/10.18653/v1/2021.emnlp-main.230>
- 1897 [116] Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. 2024. Scaling LLM Test-Time Compute Optimally Can Be More Effective Than Scaling Model Parameters. *arXiv:2408.03314 [cs.CL]*
- 1898 [117] Benjamin Spector, Andreas Kipf, Kapil Vaidya, Chi Wang, Umar Farooq Minhas, and Tim Kraska. 2021. Bounding the Last Mile: Efficient Learned String Indexing. In *AIDB*.
- 1899 [118] Ilya Sutskever, Oriol Vinyals, and Quoc V. Le. 2014. Sequence to sequence learning with neural networks. In *Proceedings of the 28th International Conference on Neural Information Processing Systems - Volume 2* (Montreal, Canada) (NIPS'14). MIT Press, Cambridge, MA, USA, 3104–3112.
- 1900 [119] Anton Tarasenko. 2017. Social-Media-Queries (SMQ). GitHub repository. <https://github.com/antontarasenko/smq/blob/master>
- 1901 [120] DeepSeek Team. 2024. DeepSeek-V3 Technical Report. <https://arxiv.org/html/2412.19437v1>
- 1902 [121] Qwen Team. 2025. Qwen-2.5 Technical Report. Technical Report 2412.15115. arXiv. <https://arxiv.org/pdf/2412.15115.pdf>

- [122] testzer0 and contributors. 2022. AMBI-QT: Ambiguous Query Tasks for NL \rightarrow SQL. GitHub repository. <https://github.com/testzer0/AmbiQT/tree/master>
- [123] Together Computer. 2023. RedPajama-Data-1T-Sample: A 1.2 B-token sample of the RedPajama pre-training corpus. <https://huggingface.co/datasets/togethercomputer/RedPajama-Data-1T-Sample> Hugging Face dataset snapshot, accessed 2025-10-24.
- [124] Ask TOM. [n.d.]. Using the LIKE predicate in a WHERE clause. <https://asktom.oracle.com/ords/asktom.search?tag=using-the-like-predicate-in-a-where-clause>. Accessed: 2025-03-07.
- [125] TPC. [n. d.]. TPC-H Homepage. <https://www.tpc.org/tph/>. Accessed: 2025-03-18.
- [126] Kapil Vaidya, Subarna Chatterjee, Eric Knorr, Michael Mitzenmacher, Stratos Idreos, and Tim Kraska. 2022. SNARF: a learning-enhanced range filter. *Proc. VLDB Endow.* 15, 8 (April 2022), 1632–1644. <https://doi.org/10.14778/3529347.3529347>
- [127] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. In *Proceedings of the 31st International Conference on Neural Information Processing Systems* (Long Beach, California, USA) (*NIPS’17*). Curran Associates Inc., Red Hook, NY, USA, 6000–6010.
- [128] Youyun Wang, Chuzhe Tang, Zhaoguo Wang, and Haibo Chen. 2020. SIndex: a scalable learned index for string keys. In *Proceedings of the 11th ACM SIGOPS Asia-Pacific Workshop on Systems* (Tsukuba, Japan) (*APSys ’20*). Association for Computing Machinery, New York, NY, USA, 17–24. <https://doi.org/10.1145/3409963.3410496>
- [129] WIKI. [n. d.]. Latest WIKI-dump(2025). <https://dumps.wikimedia.org/enwiki/>. Accessed: 2025-10-21.
- [130] Jiacheng Wu, Yong Zhang, Shimin Chen, Jin Wang, Yu Chen, and Chunxiao Xing. 2021. Updatable Learned Index with Precise Positions. *Proc. VLDB Endow.* 14, 8 (apr 2021), 1276–1288. <https://doi.org/10.14778/3457390.3457393>
- [131] Shiguang Wu, Yaqing Wang, and Quanming Yao. 2025. Why In-Context Learning Models Are Good Few-Shot Learners?. In *The Thirteenth International Conference on Learning Representations*. <https://openreview.net/forum?id=iUcsecZJp>
- [132] Ziniu Wu, Ryan Marcus, Zhengchun Liu, Parimarjan Negi, Vikram Nathan, Pascal Pfeil, Gaurav Saxena, Mohammad Rahman, Balakrishnan Narayanaswamy, and Tim Kraska. 2024. Stage: Query Execution Time Prediction in Amazon Redshift. In *Companion of the 2024 International Conference on Management of Data* (Santiago AA, Chile) (*SIGMOD/PODS ’24*). Association for Computing Machinery, New York, NY, USA, 280–294. <https://doi.org/10.1145/3626246.3653391>
- [133] Chuau Xiao, Wei Wang, Xuemin Lin, Jeffrey Xu Yu, and Guoren Wang. 2011. Efficient similarity joins for near-duplicate detection. *ACM Trans. Database Syst.* 36, 3, Article 15 (Aug. 2011), 41 pages. <https://doi.org/10.1145/2000824.2000825>
- [134] XJTU Intelligent Software Lab. 2021. CHASE: Chinese Text-to-SQL Dataset with Logical Reasoning. Project website. <https://xjtu-intsoft.github.io/chase/>
- [135] XLang-AI. 2023. Spider 2.0. GitHub repository. <https://github.com/xlang-ai/Spider2>
- [136] Xiaochun Yang, Bin Wang, and Chen Li. 2008. Cost-based variable-length-gram selection for string collections to support approximate queries efficiently. In *Proceedings of the 2008 ACM SIGMOD International Conference on Management of Data* (Vancouver, Canada) (*SIGMOD ’08*). Association for Computing Machinery, New York, NY, USA, 353–364. <https://doi.org/10.1145/1376616.1376655>
- [137] Yifan Yang and Shimin Chen. 2024. LITS: An Optimized Learned Index for Strings. *Proc. VLDB Endow.* 17, 11 (July 2024), 3415–3427. <https://doi.org/10.14778/3681954.3682010>
- [138] Zongheng Yang, Amog Kamsetty, Sifei Luan, Eric Liang, Yan Duan, Xi Chen, and Ion Stoica. 2020. NeuroCard: One Cardinality Estimator for All Tables. *Proc. VLDB Endow.* 14, 1 (sep 2020), 61–73. <https://doi.org/10.14778/3421424.3421432>
- [139] Zongheng Yang, Eric Liang, Amog Kamsetty, Chenggang Wu, Yan Duan, Xi Chen, Pieter Abbeel, Joseph M. Hellerstein, Sanjay Krishnan, and Ion Stoica. 2019. Deep Unsupervised Cardinality Estimation. *Proc. VLDB Endow.* 13, 3 (nov 2019), 279–292. <https://doi.org/10.14778/3368289.3368294>
- [140] Shaoxiong Yu, Li Han, Marta Indulska, and Hongzhi Liu. 2023. Human-in-the-loop Regular Expression Extraction for Single Column Format Inconsistency. In *WWW*, 3859–3867.
- [141] Tao Yu, Rui Zhang, Oleksandr Polozov, Christopher Meek, and Dragomir Radev. 2019. CoSQL: A Conversational Text-to-SQL Challenge Towards Cross-Domain Natural Language Interfaces to Databases. In *Proceedings of EMNLP-IJCNLP 2019*. <https://aclanthology.org/D19-1204/>
- [142] Tao Yu, Rui Zhang, Kai Yang, Michihiro Yasunaga, Dongxu Wang, Zifan Li, James Ma, Irene Li, Qingning Yao, Shanelle Roman, Zilin Zhang, and Dragomir Radev. 2018. Spider: A Large-Scale Human-Labeled Dataset for Complex and Cross-Domain Semantic Parsing and Text-to-SQL Task. In *Proceedings of EMNLP 2018*. Association for Computational Linguistics, Brussels, Belgium, 3911–3921. <https://doi.org/10.18653/v1/D18-1425>

2031
2032
2033
2034
2035
2036
2037
2038
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088

APPENDIX NAVIGATION

This appendix provides a comprehensive suite of empirical analyses, scalability studies, and implementation details that underpin the claims and contributions of our paper. To help readers navigate this rich supplementary material, we offer the following roadmap:

- (1) **Section A (An Empirical Analysis of Real-World LIKE Predicate Usage)** presents the foundational study of 84,047 real-world LIKE predicates across 30 diverse datasets. It characterizes the distribution of wildcards (%) and (_), validates the semantic nature of masked substrings, and analyzes selectivity and positional correlations—providing the empirical motivation for our approach.
- (2) **Section B (In-depth Scalability Analysis)** investigates how our method scales to billion-record datasets and long-text predicates. It demonstrates that high performance can be achieved with modest training data fractions and that a fine-tuned GPT-2 model enables efficient, high-recall prefix generation for inputs up to 8,192 characters with sub-100ms latency.
- (3) **Section C (Empirical Validation of Theoretical Guarantees)** offers experimental evidence for our theoretical framework, confirming that pattern entropy predicts query difficulty and that increasing model capacity yields a clear scaling law for recall—especially for high-entropy patterns.
- (4) **Section D (In-Depth Performance Analysis against Trigram-based Indexes)** conducts a fine-grained comparison with GIN indexes across three workload types, revealing SMILE's superiority on conjunctive low-selectivity queries and its robustness to underscore (_) fragmentation.
- (5) **Section E (Workload Construction)** details the algorithmic procedure used to generate linguistically grounded synthetic workloads (W1 and W2), ensuring realism through syllable- and morpheme-level masking.
- (6) **Section F (Optimizing Baselines via Parallelism and Partitioning)** documents our rigorous effort to maximize baseline performance through CPU parallelism and data partitioning, establishing a strong and fair performance ceiling governed by Amdahl's Law.
- (7) **Section G (Cost-Effectiveness Model)** introduce the unified cost model.
- (8) **Section H (An In-depth Study of End-to-End Performance)** validates SMILE's practical impact on full query execution in standard benchmarks (TPC-H and IMDB-JOB), demonstrating consistent latency reduction and over 95% recall while contextualizing the relationship between predicate-level acceleration and end-to-end gains.
- (9) **Section I (Theoretical results)** provides the detailed derivations for the theoretical results presented in Section 5.1.

Together, these sections provide rigorous, reproducible, and multi-dimensional validation of our system's design, efficacy, and practicality in real-world database environments.

A AN EMPIRICAL ANALYSIS OF REAL-WORLD LIKE PREDICATE USAGE

This chapter presents a comprehensive empirical study on the usage patterns of LIKE predicates in real-world scenarios. Our primary

objective is to uncover and characterize the prevalent patterns that emerge from a large-scale analysis of diverse datasets.

Our investigation is structured to answer several key research questions: (1) How does this large-scale survey differ from and extend previous work in this area? (2) What is the statistical distribution of the percentage sign ("%) wildcard in typical "LIKE" predicates? (3) Similarly, what is the distribution of the underscore ("_") wildcard? (4) What is the interplay and co-occurrence relationship between the "%" and "_" wildcards within the same predicate?

A.1 Comparison with Prior Work

While our study is the most extensive to date, we build upon foundational work that first highlighted the importance of "LIKE" predicate analysis. A notable example is the work presented in [75], which investigated patterns within the well-known IMDB Join Order Benchmark (JOB). The authors made a key observation that: "LIKE predicates appearing in the Join Order Benchmark (JOB) workload have typically 2 or 3 %s, and there is a single or two words between two contiguous %s."

However, while this seminal study provided valuable initial insights, its conclusions are drawn from the relatively narrow context of the JOB workload, which comprises only 195 "LIKE" predicates from a single domain. The benchmark's primary focus on join order optimization means it may not fully represent the diverse ways "LIKE" predicates are utilized across the broader spectrum of database applications. To establish a more general and robust understanding, our research broadens this scope significantly. We perform a large-scale analysis of 84,047 "LIKE" predicates extracted from 30 distinct and diverse datasets, offering substantial advantages in scale, diversity, and the breadth of application contexts. The details of the workloads we collected are shown in Table 2. The complete workload, including all datasets link and predicates, is publicly available at "https://anonymous.4open.science/r/SMILE-0D8B/Realworld_84047LIKE.csv". Details are summarized in Table 2.

Grounded in this extensive and diverse corpus, our study systematically dissects the characteristics of real-world "LIKE" predicates. We begin by analyzing the distribution and structure of the most common wildcards, then delve into the semantic meaning of the enclosed keywords, and finally investigate their impact on query selectivity and execution complexity. This multi-faceted analysis validates our core hypotheses and reveals key characteristics that motivate our proposed approach.


A.2 Distribution and Patterns of the % Wildcard

We begin by investigating the usage of the most common wildcard, "%". Our analysis aims to validate that patterns with leading and/or trailing wildcards, such as "%key%" and "%key1%key2%", are the most prevalent forms that challenge traditional prefix-based indexing.

Figure 16(a) illustrates the frequency of "%" occurrences within a single "LIKE" predicate. The y-axis is on a logarithmic scale to accommodate the wide distribution. It is evident that queries containing one or two "%" wildcards are overwhelmingly dominant,

Table 2: Overview of realistic SQL workloads employed in our evaluation

Dataset Source	Num of LIKEs	Avg. Num of % Signs per LIKE	Number of LIKEs Containing _	Length Distribution Of LIKEs	Category	Avg.Table.Rows
AMBI-QT [122]	1142	1.80	0	5 7 10 21	Fine-tuning Analysis	1838
ST2S [92]	3383	1.65	24	8 12 16 38	Pretraining / Text-to-SQL	2
SQL-Advising [73]	271	0.90	12	6 13 14 14	Fine-tuning Analysis	8807
BULL [1]	159	1.39	0	5 11 17.2 35	General LLM DB Reasoning	3294
CSS [2]	390	1.87	0	4 5 5 6	Chinese Domain SQL	7421
Cpiser [94]	207	0.56	0	2 3 6 18	Chinese Domain SQL	17
JOB [3]	194	2.00	0	8 10 13 35	Query Optimization	3526284
KaggleDBQA [78]	18	1.88	0	9 12.5 17.5 24	Fine-tuning Analysis	280035
EHIRSQL [79]	1479	2.00	0	8 8 8 8	Medical SQL QA	108000
LogicCat [4]	66	1.87	0	8.5 14 16 27	Logical Reasoning	20
SQL-NLP [99]	24	1.58	0	9 10 14 14	Educational Demo	24441
SocialMediaQueries [119]	21	1.95	4	10 15 19 23	Social Media SQL QA	45649087
SpiderSS [5]	188	1.81	0	5 7.25 11 21	Data Augmentation	1838
TPCH [31]	7	1.71	0	4 6.5 12.6 21	SQL Engine Benchmark	3400000(SF=1)
Advising [72]	126	1.80	0	8 13 14 14	Fine-tuning Analysis	8807
Beaver [6]	31	1.03	0	3 3 8 17	Enterprise SQL QA	46108
Bird [19]	833	0.76	7	6 10 22 178	Complex JOIN Reasoning	707332
Chase [134]	627	1.69	0	4 6 8 21	Logical Parsing	999
Cosql [141]	21	2.00	0	6 10 11 16	Conversational SQL QA	1838
Dusql [106]	119	0.00	0	2 2 3 5	Chinese SQL QA	5
Sciencebench [69]	51	1.82	0	5 7 11 19	Scientific Domain SQL QA	1256108
Sede [59]	6275	1.84	70	9 13 18 104	NL2SQL	40903150
Sesql [77]	614	1.66	5	4 4 5 8	Chinese SQL QA	5
Sparc [143]	8	2.00	0	5.5 7.75 10 10	Multi-turn SQL Generation	1838
Spider [142]	281	1.62	0	6 9 13 42	Generalization and Schema Linking	1838
Spider-DK [53]	14	1.57	0	5 5.75 6.7 7	Data Augmentation	10
Spider2 [135]	90	1.67	0	8 10 12 21	Large scale DB QA	5273M
Spidersyn [52]	176	1.80	0	5 8 11 21	Data Augmentation	1838
OmniSQL [112]	67058	1.62	1467	8 11 16 110	Pretraining / Text-to-SQL	16
Vanna [11]	174	2.16	0	12 21 43 91	Fine-tuning Analysis	98339



(a) Frequency of "%" per query (log scale).
(b) Top-5 most frequent LIKE patterns.
Figure 16: Distribution analysis of the "%" wildcard in LIKE predicates. (a) shows the number of "%" characters per query, highlighting the dominance of 1- and 2-wildcard queries. (b) shows the absolute frequency (blue) and dataset prevalence (orange) of the most common patterns.

accounting for 27,485 and 54,309 instances, respectively. This confirms that multi-keyword substring searches are a widespread phenomenon.

Figure 16(b) delves deeper into the structural patterns of these queries. The "%key%" pattern is, by a significant margin, the most frequent, appearing over 34,000 times across 29 distinct datasets. The "key%" pattern follows, with over 20,000 occurrences. Notably, patterns with multiple keywords, such as "%key1%key2%", are also highly common, ranking fifth in frequency. The dataset prevalence (orange bars) further underscores that these index-unfriendly patterns are not artifacts of a single application but are a pervasive feature across diverse database workloads.

Takeaway: The "%key%" and "%key1%key2%" patterns are one of the most common and widespread structures in real-world "LIKE" queries, confirming and highlighting the critical need for solutions that can efficiently handle leading wildcards.

A.3 Empirical Validation of Workload Generation

To substantiate the rationale behind our realism-driven workloads (W1 and W2), we conducted an empirical study to analyze the semantic nature of substrings masked by wildcards in real-world queries. The central hypothesis is that for the most challenging LIKE patterns (i.e., those with leading wildcards), the masked portions are not arbitrary character sequences but typically correspond to a small number of fundamental linguistic units. Validating this hypothesis is crucial, as it ensures our synthetic workloads accurately reflect the challenging yet meaningful search patterns encountered in production environments.

Our investigation analyzed a corpus of 84,047 production LIKE queries, focusing on the most prevalent patterns. We employed Deepseek V3 to perform a statistical analysis, classifying the content masked by the wildcard ("%") into three categories: *morphemes* (the smallest meaningful units in a language), *syllables* (units of pronunciation), and *others*. The "Others" category encompasses substrings with low semantic or phonetic significance, such as numerical sequences, punctuation, or arbitrary character fragments. It is important to note that the morpheme and syllable categories are not mutually exclusive; a substring can represent both. For instance, in a query like "%MOD%" against the string "SIGMOD",




Figure 17: Statistical analysis of wildcard-masked content in common LIKE patterns from realistic patterns. The bars show the proportion of masked substrings that can be described by morphemes, syllables, or other content types.


the masked prefix "SIG" can be interpreted both as a syllable for pronunciation and as a morpheme representing a 'special Interest Group'. Figure 17 presents the results of this analysis. The findings strongly support our hypothesis for the most index-unfriendly patterns:

- For infix patterns such as %Key% and %Key%Key%, a significant majority of the masked substrings correspond to morphemes (approx. 66% and 50%, respectively) and syllables (approx. 67% and 53%). This indicates that users frequently formulate infix searches by omitting one or two meaningful or pronounceable chunks from the beginning, end, or middle of a string.
- Similarly, for suffix queries (%Key), the masked prefixes are predominantly morphemes (65%) and syllables (72%). This pattern suggests users often know the ending of a string but are uncertain about a well-defined prefix, which might be a name, a category, or another semantic unit.
- Conversely, for prefix patterns (Key%), the "Others" category dominates (65%). This is an insightful result, suggesting that when users know the beginning of a string, the unknown suffix is often not a clean linguistic unit but rather an arbitrary truncation. However, prefix queries are efficiently handled by traditional indexes like B+-trees and are not the primary focus of our work on accelerating challenging LIKE queries.

Takeaway: Empirical analysis of 84,047 production queries reveals that wildcards in challenging infix and suffix LIKE patterns predominantly mask a small number of semantic (morphemes) or phonetic (syllables) units. This finding validates our principled, linguistics-based approach to generating realistic and challenging workloads.

A.4 Analysis of the _ Wildcard

While the "%" wildcard is dominant, the single-character wildcard, "_", also appears in real-world queries, introducing distinct challenges due to its fixed-length matching semantics.



(a) Frequency of "_" per query.
(b) Normalized length of segments delimited by "_".
Figure 18: Usage characteristics of the "_" wildcard. (a) shows that most queries use "_" sparingly. (b) reveals that segments created by "_" are predominantly short.

As shown in Figure 18(a), the absolute frequency of "_" is substantially lower than that of "%". Most queries that use this wildcard contain only one or two instances. However, its impact on query complexity is disproportionate. Figure 18(b) reveals a critical characteristic: the length distribution of substrings (segments) delimited by "_". A vast majority of these segments are extremely short. This proliferation of short, often low-selectivity segments (e.g., single characters or stop words) poses a significant challenge for optimization.

Takeaway: "_" wildcard, though less frequent than "%", fragments query patterns into very short substrings. This fragmentation can drastically increase the optimization difficulty.

A.5 Selectivity Analysis of Keyword Conjunctions

A key question is whether %Key%Key% queries commonly feature a conjunction of low-selectivity keywords that collectively yield a high-selectivity result. Such queries may be notoriously difficult for traditional systems. To investigate this, we approximated keyword selectivity by measuring word frequencies in the large-scale sample from Common Crawl text corpus with ten billion tokens [123]. This choice is motivated by the observation that most LIKE-based predicates in practice operate over relatively small datasets. If we were to estimate selectivity directly from such small datasets, we would encounter significant statistical instability: for databases with fewer than 100 tuples, the estimated selectivity of a rare keyword would be either $1/|D|$ or 0, leading to unreliable and highly variable estimates. To avoid this circular dependency and ensure consistent, statistically robust frequency estimates, we instead adopt the Common Crawl corpus—a large, publicly available, and representative text corpus—as a unified source for approximating keyword frequencies.

Table 3: Prevalence of queries where all individual keywords have low selectivity (frequency > threshold) but their conjunction has high selectivity (frequency < threshold).

Individual Keyword Selectivity Threshold (\sim)	Joint Pattern Selectivity Threshold (\sim)	Qualifying Queries	Total Queries	Proportion (%)
1e-05	1e-07	590	1009	58.47
1e-05	1e-06	687	1009	68.09
1e-04	1e-07	190	1009	18.83
1e-04	1e-06	237	1009	23.49
1e-04	1e-05	257	1009	25.47

2321
2322
2323
2324
2325
2326
2327
2328
2329
2330
2331
2332
2333
2334
2335
2336
2337
2338
2339
2340
2341
2342
2343
2344
2345
2346
2347
2348
2349
2350
2351
2352
2353
2354
2355
2356
2357
2358
2359
2360
2361
2362
2363
2364
2365
2366
2367
2368
2369
2370
2371
2372
2373
2374
2375
2376
2377
2378
2379
2380
2381
2382
2383
2384
2385
2386
2387
2388
2389
2390
2391
2392
2393
2394
2395
2396
2397
2398
2399
2400
2401
2402
2403
2404
2405
2406
2407
2408
2409
2410
2411
2412
2413
2414
2415
2416
2417
2418
2419
2420
2421
2422
2423
2424
2425
2426
2427
2428
2429
2430
2431
2432
2433
2434
2435
2436

Table 3 quantifies this phenomenon. The results are striking. For instance, when defining "low selectivity" as a keyword frequency greater than $1e-05$ and "high selectivity" as a joint pattern frequency less than $1e-07$, we find that a remarkable **58.47%** of multi-keyword queries fit this profile. This proportion rises to **68.09%** when the joint selectivity threshold is relaxed to $1e-06$.

Takeaway: A majority of real-world multi-keyword "LIKE" queries are composed of individually common terms whose intersection is highly selective. This is a common user behavior for filtering large datasets and represents a critical workload characteristics.

A.6 Positional Correlation of % and _ Wildcards

Finally, we investigate the relationship between the positions of "%" and "_" wildcards when they co-occur in a query. Understanding this relationship is important for building a realistic workload.




Figure 19: Positional correlation between "%" and "_" wildcards in co-occurring queries. The near-zero correlation coefficients (Pearson $r = 0.002$, Spearman $p = -0.117$) indicate a lack of relationship between their normalized positions.

Figure 19 plots the average normalized position of "_" against the average normalized position of "%" for each query containing both. The scatter plot reveals no discernible pattern, appearing as a random cloud of points. This visual observation is statistically confirmed by the Pearson and Spearman correlation coefficients, which are both extremely close to zero.

Takeaway: The positions of "%" and "_" wildcards within a query are close to statistically independent. This suggests that users employ them as orthogonal tools to specify different and unrelated constraints on a string's structure, rather than using them in a coordinated, patterned way. Consequently, a system claiming full LIKE support have to prepare for the simultaneous presence of "%" and "_" and optimize accordingly by treating their positional constraints as independent variables.

B IN-DEPTH SCALABILITY ANALYSIS

In this section, we conduct a more in-depth analysis of the scalability of our proposed method. We focus on two critical aspects:

- (1) How does our methodology scale to billion-scale datasets?

(2) How does our methodology support long text inputs?

To address these questions, we selected two appropriate datasets for a comprehensive evaluation.

- **RedPajama** [123]: We utilized a 1TB corpus from the RedPajama dataset. We segmented the data into words, extract a dataset of 1 billion unique string windows, each with a maximum length of 40.
- **News Room** [56]: This dataset was used for evaluating long-text performance. Containing 1,212,741 articles with average string length being 3962.

B.1 Scalability to Billion-Scale Datasets

Training models on billion-scale datasets presents significant computational and financial challenges. This section addresses a fundamental question in large-scale data modeling: how can our method be applied to datasets comprising a billion or more records in a cost-effective manner? Specifically, we investigate several key sub-questions: Can a compact model, such as SMILE, effectively learn the data distribution and converge when trained on such massive datasets? What level of generalization performance can be expected? More importantly, is it necessary to train on the entire billion-record corpus, a process that entails significant computational and financial expenditure? Or, can we achieve high performance by training on only a small fraction of the data, thereby enabling a more cost-effective scaling strategy? In this section, we conduct two targeted experiments on the billion-record RedPajama dataset to explore these questions and demonstrate the scalability and data efficiency of our approach.

Experiment 1: Generalization on a Data Subset. This experiment evaluates the generalization performance of our model when trained on a small subset of a billion-scale dataset. We sampled two disjoint 10-million-record subsets, denoted as A and B, from the full RedPajama dataset. We trained our 16MB SMILE model on subset A and evaluated its performance on a 4,096-record sample drawn from subset B. This setup ensures that the test data was strictly held-out and entirely unseen during the training phase. The results are presented in Figure 20.




Figure 20: Generalization performance on the RedPajama billion-scale dataset. The model is trained on only 1% of the data and tested on an independent, unseen subset.

As shown in the figure, the model demonstrates strong generalization capabilities, particularly for the less complex workloads W1, W2, and W3. Specifically, with an inference sample size of 512, the recall for W1, W2, and W3 exceeds 0.88, 0.90, and 0.80, respectively. This result indicates that the model can effectively capture the underlying data distribution even when trained on a small fraction (in this case, 1%) of the total dataset. Moreover, increasing the inference sample size from 64 to 512 consistently improves recall across all

workloads, highlighting the model's ability to leverage larger sample sets at inference time to refine its predictions. Conversely, the recall for the more complex workload W4 is comparatively lower. This suggests that capturing its more intricate data patterns likely necessitates a larger training set, a hypothesis we investigate in the next experiment.

Experiment 2: Determining Sufficient Data Scale. To quantify the relationship between training data volume and model performance, we conducted a focused experiment on the challenging W4 workload from the RedPajama dataset. The objective was to identify the minimum training data fraction required to achieve high performance, thereby determining a cost-effective training strategy. We trained separate models on incrementally larger subsets of the data, ranging from 0.01% to 15% of the full dataset. The results are depicted in Figure 21.




Figure 21: Impact of training data size on generalization for the W4 workload.

The results, depicted in Figure 21, reveal a clear trend: model performance improves substantially as the training data fraction increases. Training on only 0.1% of the data yields low recall (below 0.3 for a sample size of 512), although performance can still be improved by increasing the sample size during inference. As the training data proportion grows, recall steadily climbs. For workload W4, we observe that training on approximately 10% of the data achieves excellent results: recall surpasses 0.87 with an inference sample size of 64 and reaches 0.9 with a sample size of 512. This finding is significant: it demonstrates that our method not only scales to billion-record datasets but also achieves high performance cost-effectively. By training on a strategically chosen subset (e.g., 10% for complex workloads), we can avoid the prohibitive costs associated with full-corpus training while still attaining excellent generalization.

B.2 Scalability on Long Text Predicates

To rigorously evaluate our method's performance on long text predicates, a known challenge for traditional database systems, we conducted a scalability analysis using the "News Room" dataset and our most challenging workload, W4. Many commercial DBMSs impose strict length limitations on "LIKE" patterns; for instance, SQL-Server [93] has a limit of 8192 characters, while SAP HANA's is 128 [110]. We adopted the 8192-character limit as the upper bound for our tests to demonstrate our system's ability to operate well beyond typical industrial constraints.

In this experiment, we employed a fine-tuned GPT-2-Small (124M) model, as detailed in Section 5.2, to generate 40-character prefixes for given "LIKE" patterns. A key parameter in our approach is the number of generated prefix samples, which we denote as k . To understand the trade-off between accuracy and performance, we systematically varied k from 1 to 8. The inference overhead for

this process is $O(|\text{Prefix}| + k \times |\text{Output}|)$. Since the output length is a fixed-size prefix, latency is primarily influenced by the input string length ($|\text{Prefix}|$) and the number of samples k . The results, showcasing the interplay between recall, latency, input length, and sample count, are presented in Figure 22 and Figure 23.




Figure 22: Recall vs. Input String Length for a varying number of generated samples (k). Increasing k significantly improves and stabilizes recall, with $k = 4$ consistently achieving over 95% and $k = 8$ approaching near-perfect accuracy.




Figure 23: Latency vs. Input String Length for a varying number of generated samples (k). Latency scales linearly with both input length and k , but remains well under 100ms for the highly effective $k = 4$ configuration.

The results reveal two critical insights. First, as depicted in Figure 22, our method demonstrates exceptional prefix generation efficiency. While a single sample ($k = 1$) yields volatile recall, performance dramatically improves and stabilizes as k increases. Notably, with just $k = 4$ samples, the recall consistently surpasses the 95% threshold across all tested input lengths up to 8192 characters. Increasing to $k = 8$ further pushes the recall to near-perfect levels, offering a configuration for applications demanding maximum accuracy. This underscores the superior generative capability of the GPT-2 model in this task; it achieves high coverage with a remarkably small number of samples, a stark contrast to the standard SMILE(16MB) which require 64 samples to attain similar coverage in shorter text length.

Second, this high recall is achieved with minimal latency overhead, as shown in Figure 23. The latency scales linearly and predictably with both the input length and the number of samples k . While generating more samples naturally incurs a higher computational cost, the overhead remains modest and practical. For the $k = 4$ setting, which delivers excellent recall, the end-to-end latency stays comfortably below 100ms even for the longest inputs. Even in the most demanding $k = 8$ scenario, the latency remains under 180ms. This demonstrates a highly favorable trade-off, allowing practitioners to tune the parameter k to meet specific application requirements for accuracy without incurring prohibitive performance penalties.

2611
2612
2613
2614
2615
2616
2617
2618
2619
2620
2621
2622
2623
2624
2625
2626
2627
2628
2629
2630
2631
2632
2633
2634
2635
2636
2637
2638
2639
2640
2641
2642
2643
2644
2645
2646
2647
2648
2649
2650
2651
2652
2653
2654
2655
2656
2657
2658
2659
2660
2661
2662
2663
2664
2665
2666
2667
2668

Takeaway: Our method exhibits outstanding scalability for long text queries, effectively overcoming the stringent length limitations of conventional DBMSs. By leveraging a fine-tuned GPT-2 model, it achieves over 95% recall with as few as 4 generated prefix samples, maintaining low latency (under 100ms) even for inputs up to 8192 characters. This highlights a powerful and practical trade-off between near-perfect accuracy and high performance.

C EMPIRICAL VALIDATION OF THEORETICAL GUARANTEES

In Section 5.1, we established a theoretical framework that links a LIKE pattern's intrinsic complexity, quantified by its entropy $H(\mathcal{P})$, and a model's capacity, represented by its parameters $|\theta|$, to the expected query performance. This section presents a series of targeted experiments designed to empirically validate the core hypotheses derived from our theory. Specifically, we aim to verify two key predictions: (1) that query recall is inversely correlated with pattern entropy, and (2) that increasing model capacity systematically improves recall, particularly for high-entropy queries, suggesting a scaling law for neural LIKE decoding.

C.1 Impact of Pattern Entropy on Recall

Hypothesis. Our theoretical model, particularly Eq. 4, posits an inverse relationship between a pattern's entropy $H(\mathcal{P})$ and the model's success probability, which directly translates to recall. Higher entropy implies greater uncertainty in the generative process, making it inherently more challenging for the model to produce correct values. This increased uncertainty should manifest as a performance boundary, leading to lower recall for high-entropy patterns.

Experimental Setup. To test this hypothesis, we conducted an experiment on the real-world IMDB-Name dataset. We generated a workload of 2,000 queries based on the W1 template. For each query, we calculated its pattern entropy based on the data distribution. Subsequently, we partitioned the queries into fine-grained clusters based on their entropy values and measured the average recall for each cluster.




Figure 24: Negative correlation between average recall and average pattern entropy per cluster on the IMDB dataset. Each blue dot represents a cluster: its x-coordinate is the mean entropy of the cluster's patterns, and its y-coordinate is the cluster's mean recall. The red dashed line shows the linear regression fit.

Results and Analysis. Figure 24 presents the scattered results on each cluster centroid, plotting average recall as a function of pattern entropy. The empirical data points clearly demonstrate a strong negative correlation between the two variables. As pattern entropy increases, the average recall systematically decreases. The superimposed linear regression line ($\text{Recall} \approx -0.0010 \cdot H(\mathcal{P}) + 1.038$) quantitatively confirms this downward trend, providing strong evidence for our hypothesis.

This result empirically validates our theoretical assertion. Patterns with low entropy (e.g., "John%"), correspond to states in the generative FSM with low branching factors and highly skewed transition probabilities, making them easy for SMILE to decode. Conversely, high-entropy patterns (e.g., "%son%"), introduce significant uncertainty at multiple decision points within the FSM. This inherent ambiguity acts as a performance boundary for the model. Nevertheless, it is noteworthy that even for patterns with entropy exceeding 120 nats, SMILE maintains a respectable recall above 0.8, demonstrating its resilience and ability to retain fundamental decoding capabilities even for the most challenging queries.

C.2 Scaling Model Capacity for High-Entropy Patterns

Hypothesis. While high entropy poses a challenge, our theory (Eq. 4) also predicts that this challenge can be mitigated by increasing the model's capacity $|\theta|$. A larger model possesses greater capacity M to memorize and generalize the complex transition rules of the data's underlying FSM. Therefore, we hypothesize that increasing model size will enhance its decoding capability, leading to improved recall across all entropy levels, with the most significant gains observed in the high-entropy regimes.

Experimental Setup. We extended our experiment on the IMDB W1 workload to investigate the relationship between model size, pattern entropy, and recall. We trained and evaluated six variants of SMILE with different hidden dimensions, corresponding to parameter sizes ranging from 0.09MB (HiddenDim=32) to 64.2MB (HiddenDim=1024). We then measured their recall performance across six distinct entropy ranges, from low ([4.1, 29.7]) to very high ([131.9, 157.5]).




Figure 25: Impact of model size on recall across different entropy ranges. Within each entropy bracket, larger models consistently achieve higher recall. The performance advantage of larger models becomes more pronounced as entropy increases. Notably, for the highest entropy regime (Entropy > 137), maintaining recall above 0.9 requires scaling model capacity to HiddenDim=1024 (64.2MB).

Results and Analysis. Figure 25 illustrates the results of this experiment. The findings are threefold and strongly support our hypothesis.

- (1) **Consistent Improvement:** Within any given entropy bracket, recall consistently and monotonically increases

with model size. For instance, in the moderate entropy range [55.2, 80.8], the recall improves from approximately 0.12 for the 0.09MB model to over 0.95 for the 16.5MB model.

- (2) **Widening Performance Gap:** More importantly, the performance gap between smaller and larger models widens as entropy increases. For low-entropy queries ([4.1, 29.7]), even the smallest model achieves a reasonable recall of nearly 0.2, and the largest models quickly approach perfect recall. However, for the highest-entropy queries ([131.9, 157.5]), the smallest models are nearly ineffective (recall < 0.05), whereas the largest 64.2MB model maintains a robust recall of over 0.9.
- (3) **Capacity-Entropy Coupling:** Crucially, our extended experiments reveal a quantitative relationship between model complexity and workload entropy. To maintain recall above 0.9 in high-entropy regimes ($\text{Entropy} > 106$), model capacity must scale accordingly to $\text{HiddenDim}=1024$ (64.2MB). This phenomenon directly validates the core claim of Theorem 5.2: model complexity couples with workload entropy $H(W)$. This quantitative relationship provides clear guidance for system deployment, enabling appropriate model configuration based on specific workload characteristics.

These results provide compelling evidence for the second part of our theory. Larger models, with their greater parameter count $|\theta|$, are better equipped to learn the intricate character-level distributions required to navigate the high-uncertainty state transitions defined by complex LIKE patterns. This phenomenon suggests the existence of a *scaling law for neural LIKE predicate decoding*: investing more model parameters yields predictable and significant improvements in decoding accuracy, especially for the most ambiguous and difficult segments of the query workload. This validates the crucial role of model capacity, $M = \alpha|\theta|$, in our theoretical framework and demonstrates a practical path to enhancing performance by scaling the model architecture.

Takeaway. In summary, our empirical validation yields three critical insights that directly support our theoretical framework. First, we confirm that pattern entropy is a robust and accurate predictor of query difficulty, establishing a clear performance boundary dictated by the inherent uncertainty of the LIKE pattern. Second, we demonstrate that this boundary is not static—by systematically increasing the model’s parameter count, we can enhance its capacity to learn complex data distributions, thereby improving recall, especially for the most challenging high-entropy queries. Third, and most importantly, we establish a quantitative scaling relationship: to achieve certain degree of recall, model complexity is determined by workload entropy $H(\mathbb{W})$. This reveals a promising scaling law for neural LIKE predicate decoding, providing a principled approach for trading model size for query performance to meet specific application requirements.


D IN-DEPTH PERFORMANCE ANALYSIS AGAINST TRIGRAM-BASED INDEXES

To provide a more nuanced understanding of SMILE’s performance characteristics, we conduct a head-to-head comparison against a

highly optimized trigram-based GIN (Generalized Inverted Index), a standard and powerful baseline for accelerating LIKE queries in modern database systems like PostgreSQL. The objective of this analysis is to delineate the specific query scenarios where each approach excels, focusing on the interplay between the selectivity of individual keywords within a pattern and the selectivity of their combined intersection.

Workload Design. We designed three distinct query workloads (Type-A, Type-B, and Type-C) to probe the performance of both systems under controlled conditions. We define a keyword as high-selectivity if its occurrence frequency is below 0.5% and low-selectivity if its frequency is above 2%. Furthermore, to simulate the complexity of real-world patterns, based on the findings in Section A.4, we synthetically introduced a variable number of single-character wildcards (_), ranging from 0 to 7, into each pattern independently. This range is based on our analysis of 84,047 real-world queries, where seven was the maximum observed number of such wildcards.

- **Type-A (Low-Selectivity Conjunction):** Queries follow the pattern "%low_sel1%low_sel2%". Each keyword by itself is common (low selectivity), but their combined occurrence is rare (high selectivity). This represents the ideal scenario for methods that can efficiently compute intersections.
- **Type-B (High-Selectivity Dominant):** Queries follow the pattern "%high_sel1%low_sel1%". The pattern contains at least one highly selective keyword that can be used by an index to dramatically prune the search space.
- **Type-C (Single Keyword):** Queries contain only a single keyword, "%word%", which can be of either high or low selectivity. This workload evaluates the baseline performance without the advantage of intersection.



(a) Type-A: Low-Sel Conjunction. (b) Type-B: High-Sel Dominant. (c) Type-C: Single Keyword.
Figure 26: Performance comparison of SMILE and a GIN index across three query workloads with varying keyword selectivity and an increasing number of single-character "_" wildcards. Latency is shown on a linear scale.

Results and Analysis. Figure 26 presents the latency comparison between SMILE and the GIN index across the three workloads.

In the **Type-A workload** (Figure 26a), SMILE demonstrates a significant performance advantage, outperforming the GIN index by orders of magnitude. This is because the GIN index must first retrieve the large posting lists for both low-selectivity keywords and then perform a costly intersection operation. In contrast, SMILE holistically decodes the entire pattern, directly generating candidates that satisfy the conjunctive constraint without materializing large intermediate result sets.

For the **Type-B workload** (Figure 26b), the performance of SMILE and the GIN index is largely comparable. The GIN index can

effectively leverage the high-selectivity keyword to immediately filter the dataset down to a very small number of candidates, making its strategy highly efficient. SMILE's generative approach remains competitive, but its relative advantage is diminished when a simple and highly effective index path is available to the baseline.

In the **Type-C workload** (Figure 26c), the GIN index performs on par with or slightly better than SMILE. For a single-keyword search, the GIN index reduces to a direct lookup, which is its most efficient operational mode. SMILE's generative process, while fast, incurs a baseline overhead that makes it less suitable for these simple patterns where its "intersection" advantage is nullified.

A crucial observation across all workloads is the **impact of the _ wildcard**. As the number of _ wildcards increases, the performance of the GIN index degrades noticeably. This is because each _ can break up the trigrams that the index relies on. For example, a search for "SI_M_D" cannot use trigrams like "SIG" or "MOD". The index must fall back to a less efficient strategy of scanning for all possible trigram combinations, leading to higher latency. Conversely, SMILE's performance remains remarkably stable. As a generative model, it treats _ as just another token in the sequence to be generated. Its latency is therefore insensitive to the number of such wildcards, showcasing its robustness to complex and fragmented patterns.

Takeaway. This detailed analysis reveals a clear performance trade-off between SMILE and traditional trigram-based indexes. SMILE's strength lies in its ability to holistically interpret complex patterns, making it exceptionally efficient for queries involving a conjunction of low-selectivity keywords. While its performance is competitive for queries dominated by a high-selectivity term, traditional indexes are slightly more efficient for simple, single-keyword lookups. Critically, SMILE exhibits superior robustness to patterns fragmented by single-character wildcards (_), a scenario where the structural assumptions of trigram indexes break down. This suggests that SMILE is not a universal replacement but a powerful complementary engine, particularly suited for workloads rich in complex, conjunctive, and fragmented search patterns.

E WORKLOAD CONSTRUCTION

In the following section of this chapter, we describe the process of creating our loads W1 and W2. Our pseudocode can be found in Algorithm 2:

The algorithm in Algorithm 2 proceeds as follows. In line 2, the input string is decomposed into its minimal semantic units: syllables for W_1 and morphemes for W_2 . We use the Pyphen [10] to segment the syllables and morfessor [9] to segment the morphemes. For instance, the string "SIGMOD CONFERENCE" is split into $S = ["SIG", "MOD ", "CON", "FER", "EN", "CE"]$, and $n = |S|$ (e.g., $n = 6$) is computed. In line 4, a binary random decision determines whether to generate a single-field (%Keyword%) or dual-field (%Key%key%) fuzzy pattern, thereby diversifying the structural complexity of the generated queries.

In the single-field branch (lines 5–15), the algorithm samples two integers $f, b \in \{0, 1, 2\}$ to specify how many segments at the beginning and end of S are to be obscured with the multi-character wildcard %. To ensure genuine fuzziness, if $f = b = 0$, one of

Algorithm 2 Fuzzy LIKE Pattern Generation

```

Input: Text  $T$ , Options  $O$ 
Output: Fuzzy LIKE pattern  $P$ 
procedure GENERATEFUZZYPATTERN( $T, O$ )
  2:    $S \leftarrow \text{Segment}(T, O)$             $\triangleright$  Extract syllables/morphemes
  3:    $n \leftarrow |S|$ 
  4:   if  $\text{rand}() < 0.5$  then
  5:     /* Single-field pattern */
  6:      $f \leftarrow \text{rand}(0, 2); b \leftarrow \text{rand}(0, 2)$ 
  7:     if  $f + b = 0$  then
  8:        $f \leftarrow 1$  or  $b \leftarrow 1$             $\triangleright$  Ensure fuzziness
  9:        $P \leftarrow "%"$ 
 10:    for  $i = 0$  to  $n - 1$  do
 11:      if  $i < f$  or  $i \geq n - b$  then
 12:         $P \leftarrow P + "%"$ 
 13:      else
 14:         $P \leftarrow P + S[i]$ 
 15:     $P \leftarrow P + "%"$ 
 16:  else
 17:    /* Dual-field pattern */
 18:     $f \leftarrow \text{rand}(0, 2); m \leftarrow \text{rand}(0, 2); b \leftarrow \text{rand}(0, 2)$ 
 19:    if  $f + m + b = 0$  then
 20:       $f \leftarrow 1$  or  $m \leftarrow 1$  or  $b \leftarrow 1$ 
 21:    if  $m > 0$  then
 22:       $mid\_pos \leftarrow \text{rand}(f, n - b - 2)$     $\triangleright$  Middle position
 23:       $P \leftarrow "%"$ 
 24:    for  $i = 0$  to  $n - 1$  do
 25:      if  $i < f$  or  $i \geq n - b$  or ( $m > 0$  and  $i = mid\_pos$ )
 26:        then
 27:           $P \leftarrow P + "%"$ 
 28:        else
 29:           $P \leftarrow P + S[i]$ 
 30:     $P \leftarrow \text{ReplaceRandomChars}(P, "_", \text{rand}(0, 5))$     $\triangleright$  Add '_s
 31:    return  $P$ 

```

them is set to 1. The pattern P is initialized with a leading %, then constructed by iterating over all segments: those in the first f or last b positions are replaced by %, while the intervening segments are appended verbatim. A trailing % is appended to complete the pattern. For example, with $f = 1$ and $b = 2$, the resulting pattern is %MOD CONFER%, which matches any string containing the exact substring "MOD CONFER".

In the dual-field branch (lines 17–29), three integers $f, m, b \in \{0, 1, 2\}$ control fuzziness at the front, a potential middle gap, and the back, with the constraint that at least one of them is non-zero. If $m > 0$, a valid middle position mid_pos is randomly selected from the range $[f, n - b - 2]$ to avoid overlap with the front and back regions. During pattern construction, wildcards are inserted at the front, back, and—if $m > 0$ —at mid_pos . For example, with $f = 1$, $m = 1$, $b = 1$, and $mid_pos = 3$ (corresponding to the segment "FER"), the resulting pattern is %MOD CON%EN%, which requires the non-contiguous substrings "MOD CON" and "EN" to appear in order, separated by arbitrary characters.

Finally, in line 30, the algorithm invokes ReplaceRandomChars to substitute up to six randomly selected characters in the constructed pattern P with the single-character wildcard $_$. This independent random underscore replacement step simulates the pattern of load occurrence in real-world scenarios.

F OPTIMIZING BASELINES VIA PARALLELISM AND PARTITIONING

In this section, we detail our efforts to maximize the performance of the baseline methods through the strategic use of CPU parallelism and data partitioning. Our goal is to establish a robust performance benchmark, ensuring that each baseline is configured to its optimal potential before comparison. We utilize the W3 on TPC-H lineitem table as a representative workload to illustrate the outcomes of our extensive tuning process.

The core of our optimization strategy revolves around two primary axes: scaling the degree of parallelism by increasing the number of parallel workers and varying the data fragmentation by adjusting the number of partitions. Figure 27 presents the end-to-end latency of five baseline methods—GIN, GIST, Prefix B⁺-tree, Sequential Scan, and Suffix B⁺-tree—as a function of the number of parallel workers, across four distinct partitioning schemes (P=4, 8, 16, and 32).

A key observation across all subfigures is the initial sharp decline in latency as the number of parallel workers increases from one. This demonstrates the inherent parallelizability of the tasks and the effectiveness of distributing the workload across multiple CPU cores. However, this performance gain is not linear and quickly reaches a point of diminishing returns. For instance, in Figure 27(a) with 4 partitions, the latency for most baselines stabilizes after approximately 16 parallel workers. Beyond this threshold, the addition of more workers yields negligible performance improvements and, in some cases, introduces minor performance jitter.

This phenomenon is a classic illustration of Amdahl's Law [12], which posits that the speedup of a program using multiple processors is limited by its sequential fraction. In our context, factors such as I/O bottlenecks, data synchronization overhead, and the inherent sequential components of the query processing algorithms constrain the achievable speedup. Despite our best efforts to parallelize every possible aspect of the baselines, these fundamental limitations dictate a ceiling on performance scalability.

Furthermore, we systematically increased the number of data partitions from 4 to 32, as shown in subfigures (a) through (d), to further enhance the potential for parallel execution. Data partitioning can significantly improve query performance by dividing a large dataset into smaller, more manageable chunks, allowing queries to run in parallel across these segments. As depicted, increasing the partition count generally lowers the overall latency floor for the index-based methods. For example, the latency for GIN and GIST methods at 32 partitions (Figure 27(d)) is visibly lower than at 4 partitions (Figure 27(a)) when using a high degree of parallelism. This is because a higher partition count allows for a finer-grained distribution of work, reducing contention and improving load balancing among the parallel workers.

However, even with 32 partitions, the performance of all baselines eventually plateaus, reinforcing the constraints imposed by

Amdahl's Law. The sequential scan, being the most fundamentally I/O-bound operation, benefits the least from increased parallelism and partitioning, exhibiting a relatively flat performance curve after an initial small dip. This comprehensive tuning process confirms that we have pushed the conventional baselines to their practical performance limits under this parallel execution model. The observed performance saturation underscores the necessity for a more advanced approach to overcome these inherent scalability challenges.

Takeaway Our exhaustive tuning demonstrates that the performance of traditional methods is fundamentally constrained by Amdahl's Law. Despite aggressive parallelism and data partitioning, their scalability quickly reaches a point of diminishing returns due to sequential bottlenecks. This establishes a practical performance ceiling, confirming that merely adding more hardware resources is an insufficient strategy and motivating the need for architecturally superior solutions.

G DETAILED COST MODEL

Our cost estimation is based on prevailing cloud computing rental prices, which reflect the real-world deployment costs of database systems. We established two primary hardware configurations for our experiments, with rental rates sourced from public cloud providers [13, 30] and detailed in Table 4.

A critical aspect of our cost model is the selection of CPU core counts to ensure a fair and cost-effective comparison. As demonstrated in our analysis in Section F, the performance of all traditional baselines (B+-Tree, GIN, GIST, and SeqScan) saturates when the degree of parallelism reaches approximately 16 workers. Allocating additional CPU cores beyond this threshold yields negligible latency improvements due to inherent bottlenecks, as dictated by Amdahl's Law. Consequently, to model a realistic and economically rational deployment, we benchmark the traditional methods on a 16-core CPU configuration. This represents the most cost-effective setup, as provisioning the full 96-core CPU would not enhance their performance but would unfairly inflate their cost. Our method and other deep learning-based baselines, which are designed for such hardware, utilize the CPU+GPU configuration.

Table 4: Hardware Configurations and Hourly Rental Costs.

Hardware Configuration	Hourly Rental Cost (\$)
1 × AMD EPYC 9654 CPU (96 Cores)	\$1.50
1 × AMD EPYC 9654 CPU (16 Cores)	\$0.25
1 × NVIDIA H20 GPU + 1 × AMD EPYC 9654 CPU (16 Cores)	\$0.98

We calculate the cost to process one million queries, denoted as C_{1M} , based on the total execution time. Specifically, if a method takes T_{1M} seconds to execute one million queries on a hardware configuration with an hourly rate of R , the cost is computed as:

$$C_{1M} = \frac{T_{1M}}{3600} \times R$$

This metric provides a normalized and direct comparison of the economic efficiency of each method.

3133 H AN IN-DEPTH STUDY OF END-TO-END 3134 PERFORMANCE

3135 In this section, we conduct a comprehensive analysis of the end-to-
3136 end (E2E) performance of SMILE on complex queries, as extended in
3137 Sec. 5.2. The primary objective is to quantify its practical impact on
3138 full query execution and to validate its efficacy against established
3139 indexing techniques across disparate data domains. To this end, we
3140 perform experiments on two standard benchmarks: TPC-H [125]
3141 and IMDB-JOB [3]. We compare SMILE against a sequential scan
3142 baseline and several state-of-the-art index structures, including GIN,
3143 GIST, and B+ Tree variants tailored for prefix and suffix matching.
3144

3145 Specifically, we selected five representative query templates
3146 from TPC-H (Q2, Q9, Q13, Q14, Q20) and five from IMDB-JOB
3147 (Q10, Q15, Q19, Q23, Q29). These queries were chosen for their
3148 reliance on complex LIKE predicates. The TPC-H queries target
3149 the part.name, part.type, and order.comment columns (abbreviated as
3150 partName, partType, and orderComment). The IMDB-JOB queries focus on the cast_info.note and movie_info.info
3151 columns (abbreviated as castInfo and movieInfo). We measure
3152 E2E latency and conduct a granular analysis of SMILE's core filtering
3153 efficacy and predicate evaluation speed.

3154 Figure 29 presents the E2E query latency results. As illustrated,
3155 SMILE consistently and substantially outperforms all baselines
3156 across all tested queries on both benchmarks. The performance
3157 gap is particularly pronounced on complex queries like TPC-H
3158 Q9 and IMDB-JOB Q13, where SMILE achieves up to an order of
3159 magnitude reduction in latency compared to the traditional index-
3160 based approach. These results confirm the practical utility and
3161 superiority of SMILE in reducing overall query execution time. The
3162 consistent improvements on both the structured data of TPC-H and
3163 the varied textual data in IMDB-JOB underscore the robustness and
3164 generalizability of our learned approach.

3165 While these E2E results are compelling, a more granular analysis
3166 is necessary to isolate the core contribution of our method. The
3167 true potency of SMILE lies in its ability to drastically accelerate the
3168 evaluation of the LIKE predicate itself, which is often a primary bot-
3169 tleneck. To quantify this effect, we conducted a micro-benchmark
3170 to measure the predicate evaluation latency and the corresponding
3171 recall rate.

3172 Figure 28 provides this detailed breakdown. Figure 28a isolates
3173 the performance of LIKE predicate evaluation, revealing that SMILE
3174 is faster than all baselines by a significant margin. Compared to a
3175 sequential scan, SMILE achieves an acceleration of up to two orders
3176 of magnitude across all targeted columns. For instance, on the TPC-
3177 H partType and IMDB-JOB castInfo columns, SMILE's latency is
3178 orders of magnitude lower than that of SeqScan and substantially
3179 better than other specialized indexes like GIN. This demonstrates
3180 SMILE's efficacy as a specialized accelerator for pattern matching.
3181 Concurrently, Figure 28b confirms that this remarkable speedup
3182 does not come at a great cost to accuracy. SMILE maintains a high
3183 recall rate, typically exceeding 90% for most TPC-H queries. While
3184 it demonstrates some variability on the more heterogeneous IMDB-
3185 JOB dataset (e.g., JOB-Q15), its recall remains effective for pre-
3186 filtering, ensuring that the vast majority of qualifying rows are
3187 correctly identified.

3191 The stark contrast between the multi-order-of-magnitude speedup
3192 at the predicate level (Figure 28a) and the more modest, albeit still
3193 substantial, gains in E2E queries (Figure 29) is an important and
3194 expected finding. SMILE functions as a highly specialized acceler-
3195 ator for LIKE predicates, analogous to how a B-tree or learned
3196 index accelerates point/range lookups. By transforming a linear-
3197 time scan for pattern matching into a near-constant-time operation,
3198 it removes a critical bottleneck. However, in a complex analytical
3199 query, this predicate evaluation is just one component of a larger
3200 execution plan. The total query time is still heavily influenced by
3201 other substantial costs, such as I/O for data retrieval, CPU cycles
3202 for join processing, and final aggregations. These factors, which
3203 are orthogonal to our contribution, can dominate the overall exec-
3204 ution time and thus dilute the exponential gains achieved at the
3205 single-predicate level. Nonetheless, by fundamentally accelerating
3206 the LIKE operator, SMILE delivers tangible and crucial value to the
3207 performance of modern database systems.

Takeaway: SMILE delivers consistent and practical end-to-end query performance improvements, significantly outperforming both sequential scans and traditional index structures across diverse datasets. Its core strength lies in accelerating single-column LIKE predicates by up to two orders of magnitude over sequential scans while maintaining high recall (often >90%), making it a potent pre-filter. The dilution of this effect in E2E queries is an expected outcome attributable to other system bottlenecks (e.g., I/O, joins), but the fundamental acceleration of the pattern-matching operator remains a significant and valuable contribution.

I PROOFS OF THEORETICAL RESULTS

This appendix provides the detailed derivations for the theoretical results presented in Section 5.1.

I.1 Proof of Theorem 5.1

PROOF. The theorem aims to compute the pattern entropy $H(\mathcal{P}) = \ln |L(\mathcal{P})|$, where $|L(\mathcal{P})|$ is the cardinality of the language $L(\mathcal{P})$ generated by a LIKE pattern \mathcal{P} . The proof proceeds by enumerating the number of unique strings that can match the pattern.

A LIKE pattern \mathcal{P} of length x consists of literal characters, p underscore wildcards ("_"), and q percent wildcards ("%"). The alphabet is Σ , and the maximum length of a generated string is m . The cardinality $|L(\mathcal{P})|$ is the product of the number of combinations for each component of the pattern.

- (1) **Literal Characters:** The non-wildcard characters are fixed and each contributes a multiplicative factor of 1 to the total cardinality.
- (2) **Underscore Wildcards ("_"):** Each of the p underscore wildcards matches exactly one character from the alphabet Σ . As these choices are independent, they collectively contribute a factor of $|\Sigma|^p$.

$$C_{_} = |\Sigma|^p \quad (6)$$

- (3) **Percent Wildcards ("%"):** Each of the q percent wildcards can match a sequence of characters of any length, including zero. Let s be the total number of characters matched by all q percent wildcards combined. The length of a generated

3249 string is determined by the length of the fixed parts of the
 3250 pattern plus the total length s . The fixed part (literals and
 3251 underscores) has a length of $x - q$. Thus, the total string
 3252 length is $(x - q) + s$. Given the maximum length constraint
 3253 m , we have $(x - q) + s \leq m$, which implies that $s \leq m - x + q$.
 3254 The minimum value for s is 0.

3255 For a fixed total length s , we determine the number of ways
 3256 these s characters can be generated. This involves two steps:
 3257 first, we partition the total length s among the q wildcards.
 3258 This is a classic combinatorial problem solvable with stars
 3259 and bars, yielding $\binom{s+q-1}{q-1}$ ways to distribute the lengths.
 3260 Second, for each such partition, the s character positions
 3261 can be filled with any character from Σ , giving $|\Sigma|^s$ possible
 3262 sequences.

3263 To find the total contribution from the percent wildcards,
 3264 we sum over all possible values of s :

$$C_{\%} = \sum_{s=0}^{m-x+q} \binom{s+q-1}{q-1} |\Sigma|^s \quad (7)$$

3265 The total cardinality of the language $L(\mathcal{P})$ is the product of the
 3266 contributions from the underscore and percent wildcards:

$$|L(\mathcal{P})| = C_{_} \cdot C_{\%} = |\Sigma|^p \left(\sum_{s=0}^{m-x+q} \binom{s+q-1}{q-1} |\Sigma|^s \right) \quad (8)$$

3267 Finally, the pattern entropy $H(\mathcal{P})$ is the natural logarithm of this
 3268 cardinality. Using the logarithm property $\ln(ab) = \ln(a) + \ln(b)$,
 3269 we arrive at the final expression:

$$\begin{aligned} H(\mathcal{P}) &= \ln(|\Sigma|^p) + \ln \left(\sum_{s=0}^{m-x+q} \binom{s+q-1}{q-1} |\Sigma|^s \right) \\ &= p \ln |\Sigma| + \ln \left(\sum_{s=0}^{m-x+q} \binom{s+q-1}{q-1} |\Sigma|^s \right) \end{aligned} \quad (9)$$

3270 This completes the proof. \square

I.2 Proof of Theorem 5.2

3271 PROOF. The theorem provides the minimum number of samples,
 3272 n , required to retrieve at least K matching strings with a confidence
 3273 of $1 - \delta$. We model the sampling process as a sequence of n inde-
 3274 pendent Bernoulli trials. Let $X_i = 1$ if the i -th sample is a correct
 3275 match (with probability p) and $X_i = 0$ otherwise. The total number
 3276 of successes is $S_n = \sum_{i=1}^n X_i$, which follows a Binomial distribution,
 3277 $S_n \sim \text{Binomial}(n, p)$, with expectation $\mathbb{E}[S_n] = np$.

3278 Our goal is to find the minimum n such that $P(S_n \geq K) \geq 1 - \delta$,
 3279 which is equivalent to bounding the lower tail probability: $P(S_n <
 3280 K) \leq \delta$. We apply a standard multiplicative Chernoff bound, looking
 3281 for the smallest n satisfying $np > K$, which for any $K < \mathbb{E}[S_n]$
 3282 states:

$$P(S_n \leq K) \leq \exp \left(-\frac{(\mathbb{E}[S_n] - K)^2}{2\mathbb{E}[S_n]} \right) \quad (10)$$

3283 To satisfy the confidence requirement, we set the right-hand side to
 3284 be at most δ , which after taking logarithms and rearranging gives
 3285 $(np - K)^2 \geq 2np \ln(1/\delta)$. We take the positive square root to find

3286 a lower bound on the expected number of successes, leads to:

$$np \geq K + \ln \left(\frac{1}{\delta} \right) + \sqrt{\ln \left(\frac{1}{\delta} \right) \left(2K + \ln \left(\frac{1}{\delta} \right) \right)} \quad (11)$$

3287 We now derive the bound on n by considering the two cases
 3288 for the success probability p , which is lower-bounded by $p \geq c \cdot$
 3289 $\min \left(1, \frac{\alpha|\theta|}{H(\mathbb{W})} \right)$ from Eq. (4).

3290 **Case 1: Model capacity is sufficient ($H(\mathbb{W}) \leq \alpha|\theta|$)**. In this
 3291 regime, the model's capacity $M = \alpha|\theta|$ meets or exceeds the work-
 3292 load entropy. We substituting $p \geq c$ into Eq. (11):

$$n \cdot c \geq K + \ln \left(\frac{1}{\delta} \right) + \sqrt{\ln \left(\frac{1}{\delta} \right) \left(2K + \ln \left(\frac{1}{\delta} \right) \right)} \quad (12)$$

3293 This directly yields the bound for the first case.

3294 **Case 2: Model capacity is limited ($H(\mathbb{W}) > \alpha|\theta|$)**. Here, the
 3295 workload entropy exceeds the model's capacity, and the success
 3296 probability is bounded by $p \geq \frac{\alpha|\theta|}{H(\mathbb{W})}$. To guarantee that the condition
 3297 in Eq. (11) is met, we must provision for the worst-case (lowest)
 3298 possible success probability. We therefore substitute $p = \frac{c \cdot \alpha|\theta|}{H(\mathbb{W})}$:

$$n \left(\frac{c \cdot \alpha|\theta|}{H(\mathbb{W})} \right) \geq K + \ln \left(\frac{1}{\delta} \right) + \sqrt{\ln \left(\frac{1}{\delta} \right) \left(2K + \ln \left(\frac{1}{\delta} \right) \right)} \quad (13)$$


3299 Solving for n gives the bound for the second case:

$$n \geq \frac{H(\mathbb{W})}{c \cdot \alpha|\theta|} \left(K + \ln \left(\frac{1}{\delta} \right) + \sqrt{\ln \left(\frac{1}{\delta} \right) \left(2K + \ln \left(\frac{1}{\delta} \right) \right)} \right) \quad (14)$$


3300 Combining both cases completes the proof of the theorem. \square

Received July 2025; revised October 2025; accepted November 2025

3365



(a) End-to-End Latency with 4 Partitions.

3379
3380
3381
3382
3383
3384
3385
3386
3387
3388
3389
3390
3391
3392
3393
3394
3395
3396
3397
3398
3399
3400
3401
3402
3403
3404
3405
3406
3407
3408
3409
3410
3411
3412
3413
3414
3415
3416
3417
3418
3419
3420
3421
3422

(b) End-to-End Latency with 8 Partitions.

3431
3432
3433
3434
3435
3436
3437
3438
3439
3440
3441
3442
3443
3444
3445
3446
3447
3448
3449
3450
3451
3452
3453
3454
3455
3456
3457
3458
3459
3460
3461
3462
3463
3464
3465
3466
3467
3468
3469
3470
3471
3472
3473
3474
3475
3476
3477
3478
3479
3480"/>

Partition = 16


(c) End-to-End Latency with 16 Partitions.

3470
3471
3472
3473
3474
3475
3476
3477
3478
3479
3480

Partition = 32


(d) End-to-End Latency with 32 Partitions.

Figure 27: End-to-end latency as a function of the number of parallel workers for different partitioning schemes on the TPC-H lineitem table. The performance of all baselines plateaus after a certain number of workers, illustrating the limits of parallelization as described by Amdahl's Law.




(a) Predicate evaluation latency per column.

Figure 28: A detailed breakdown of SMILE’s performance against baselines. (a) The latency (log-scale) of LIKE predicate evaluation on individual columns. (b) The recall rate achieved by SMILE’s filtering mechanism for each query template. Together, these results show that SMILE achieves orders-of-magnitude acceleration over sequential scans while maintaining high recall.




(b) Recall rate per query template.

Figure 28: A detailed breakdown of SMILE’s performance against baselines. (a) The latency (log-scale) of LIKE predicate evaluation on individual columns. (b) The recall rate achieved by SMILE’s filtering mechanism for each query template. Together, these results show that SMILE achieves orders-of-magnitude acceleration over sequential scans while maintaining high recall.



(a) E2E query latency on TPC-H (Queries: 2, 9, 13, 14, 20).

Figure 29: End-to-end query performance of SMILE against baselines. The figures show total execution time (log-scale) for selected queries on TPC-H and IMDB-JOB. SMILE consistently outperforms not only the sequential scan but also established index structures like GIN, GIST, and B+ Tree variants, demonstrating its superior effectiveness in accelerating real-world analytical queries.



(b) E2E query latency on IMDB-JOB (Queries: 10, 15, 19, 23, 29).

Figure 29: End-to-end query performance of SMILE against baselines. The figures show total execution time (log-scale) for selected queries on TPC-H and IMDB-JOB. SMILE consistently outperforms not only the sequential scan but also established index structures like GIN, GIST, and B+ Tree variants, demonstrating its superior effectiveness in accelerating real-world analytical queries.

3505
3506
3507
3508
3509
3510
3511
3512
3513
3514
3515
3516
3517
3518
3519
3520
3521
3522
3523
3524
3525
3526
3527
3528
3529
3530
3531
3532
3533
3534
3535
3536
3537
3538

3539
3540
3541
3542
3543
3544
3545
3546
3547
3548
3549
3550
3551
3552
3553
3554
3555
3556
3557
3558
3559
3560
3561
3562
3563
3564
3565
3566
3567
3568
3569
3570
3571
3572
3573
3574
3575
3576
3577
3578
3579
3580
3581
3582
3583
3584
3585
3586
3587
3588
3589
3590
3591
3592
3593
3594
3595
3596